



```
subroutine gpu(  
    integer ex,   
    integer n, float* out)  
    ! Compute absolute (i,j) index  
    ! of current GPU thread using  
    ! blockIdx.x * BLOCK_LENGTH  
    ! blockIdx.y * BLOCK_HEIGHT  
    ! Compute one data point  
    ! for the given  
    ! = 0.1f *   
    ! 2.0f * 
```

**KernelGen**

naïve GPU kernels generation  
from Fortran source code

Dmitry Mikushin

# Contents

- Motivation and target
- Assembling our own toolchain: schemes and details
- Toolchain usecase: sincos example
- Development schedule

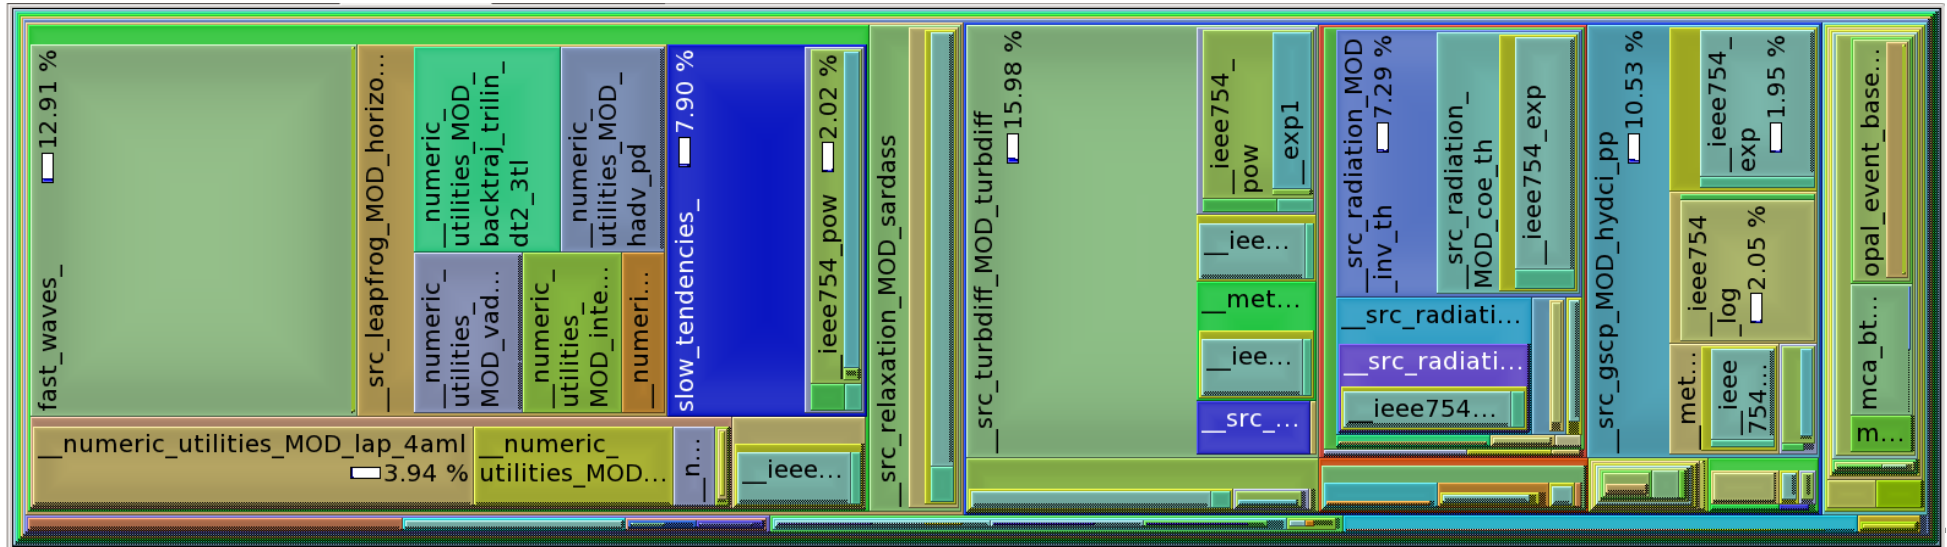
# **1. Motivation and target**



# Why generation?

The need of huge numerical models porting onto GPUs:

- All individual model blocks have too small self perf impact ( $\sim 10\%$ ), resulting into small speedups, if only one block is ported



# Why generation?

The need of huge numerical models porting onto GPUs:

- A lot of code requiring lots of similar transformations
- A lot of code versions with minor differences, each requiring manual testing & support
- COSMO, Meteo-France: science teams are not ready to work with new paradigms (moreover, tied with propriety products), compute teams have no resources to support a lot of new code

# Why generation?

So, in fact science groups are ready to start GPU-based modeling, if three main requirements are met:

- Model works on GPUs without specific extensions
- Model works on GPUs and gives accurate enough results in comparison with control host version
- Model works on GPUs faster

# Our target

Port already parallel models in Fortran onto GPUs:

- Conserving original Fortran source code (i.e. keeping all C/CUDA/OpenCL in intermediate files)
- Minimizing manual work on specific code (i.e. developed toolchain is expected to be reusable with other codes)

“Already parallel” means the model gives us some data decomposition grid to map 1 GPU onto 1 MPI process or thread.

# Similar tools

- PGI CUDA Fortran, Accelerator
- (Open)HMPP by CAPS and Pathscale
- f2c-acc

Common weaknesses: manual coding, proprietary, non-standard, non-free, closed source, non-customizable, etc.

Although, pros & cons of these toolchains is a long discussion omitted here.



## **2. Assembling our own toolchain**


# Ingredients

- **Compiler** – split original code into host and device parts and compile them into single object
  - Code splitter (source-to-source preprocessor)
  - Target device code generator
- **Runtime library** – implementation of specific internal functions used in generated code
  - Data management
  - Kernel invocation
  - Kernel results verification

# Priorities

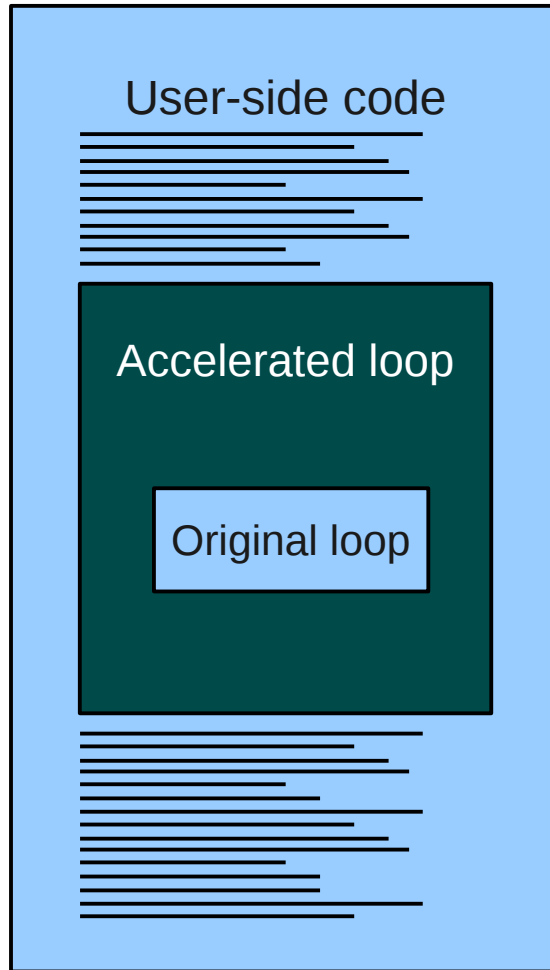
- 1) Make up the rough version of the full toolchain first, focus on improvements later
- 2) Use empirical tests where analysis is not yet sufficient (e.g. for identifying parallel loops)
- 3) Focus on best compiled kernels yield (code coverage) for COSMO and other models
- 4) Implement optimizations later

# Runtime workflow



We start with original source code, selecting loops suitable for device acceleration.

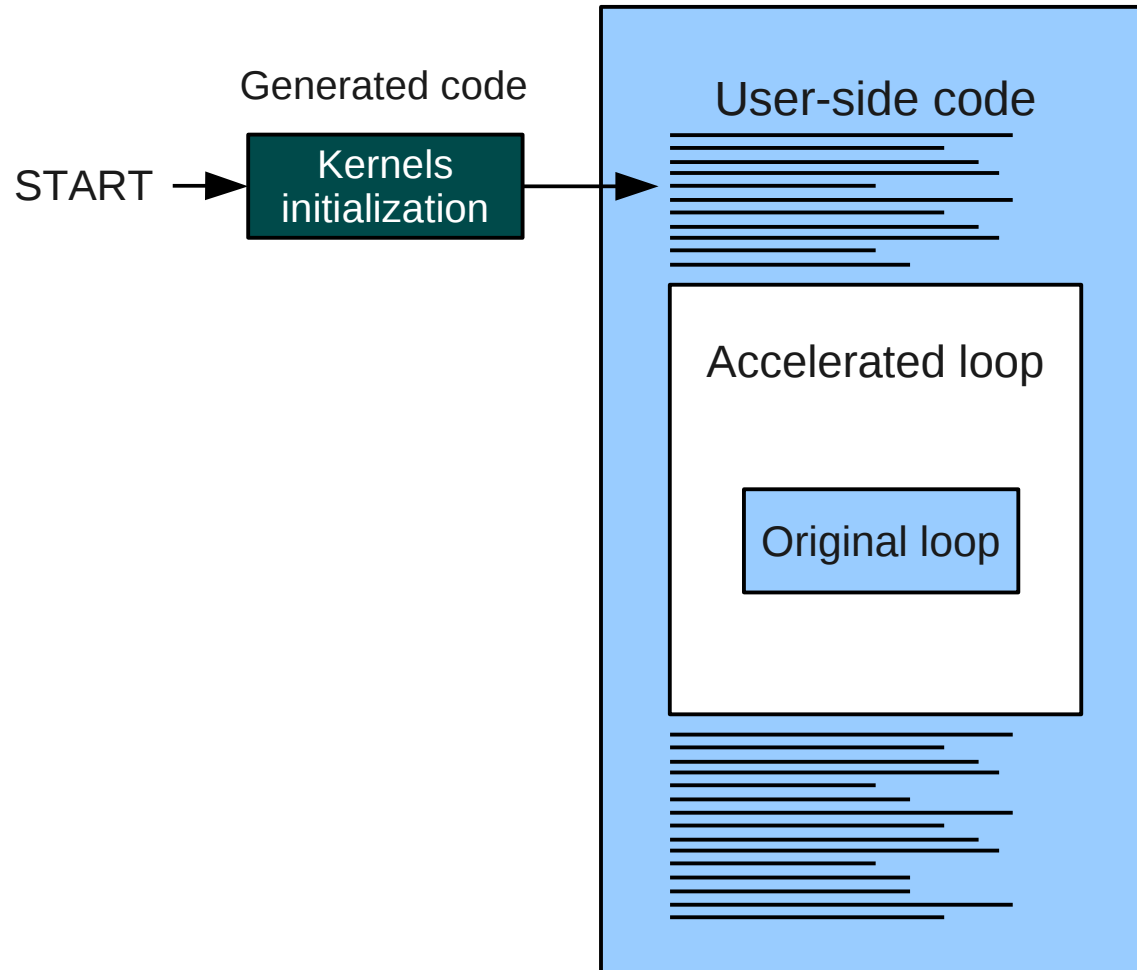
# Runtime workflow



Equivalent device code is generated for suitable loops.

(see “Code generation workflow” for details)

# Runtime workflow

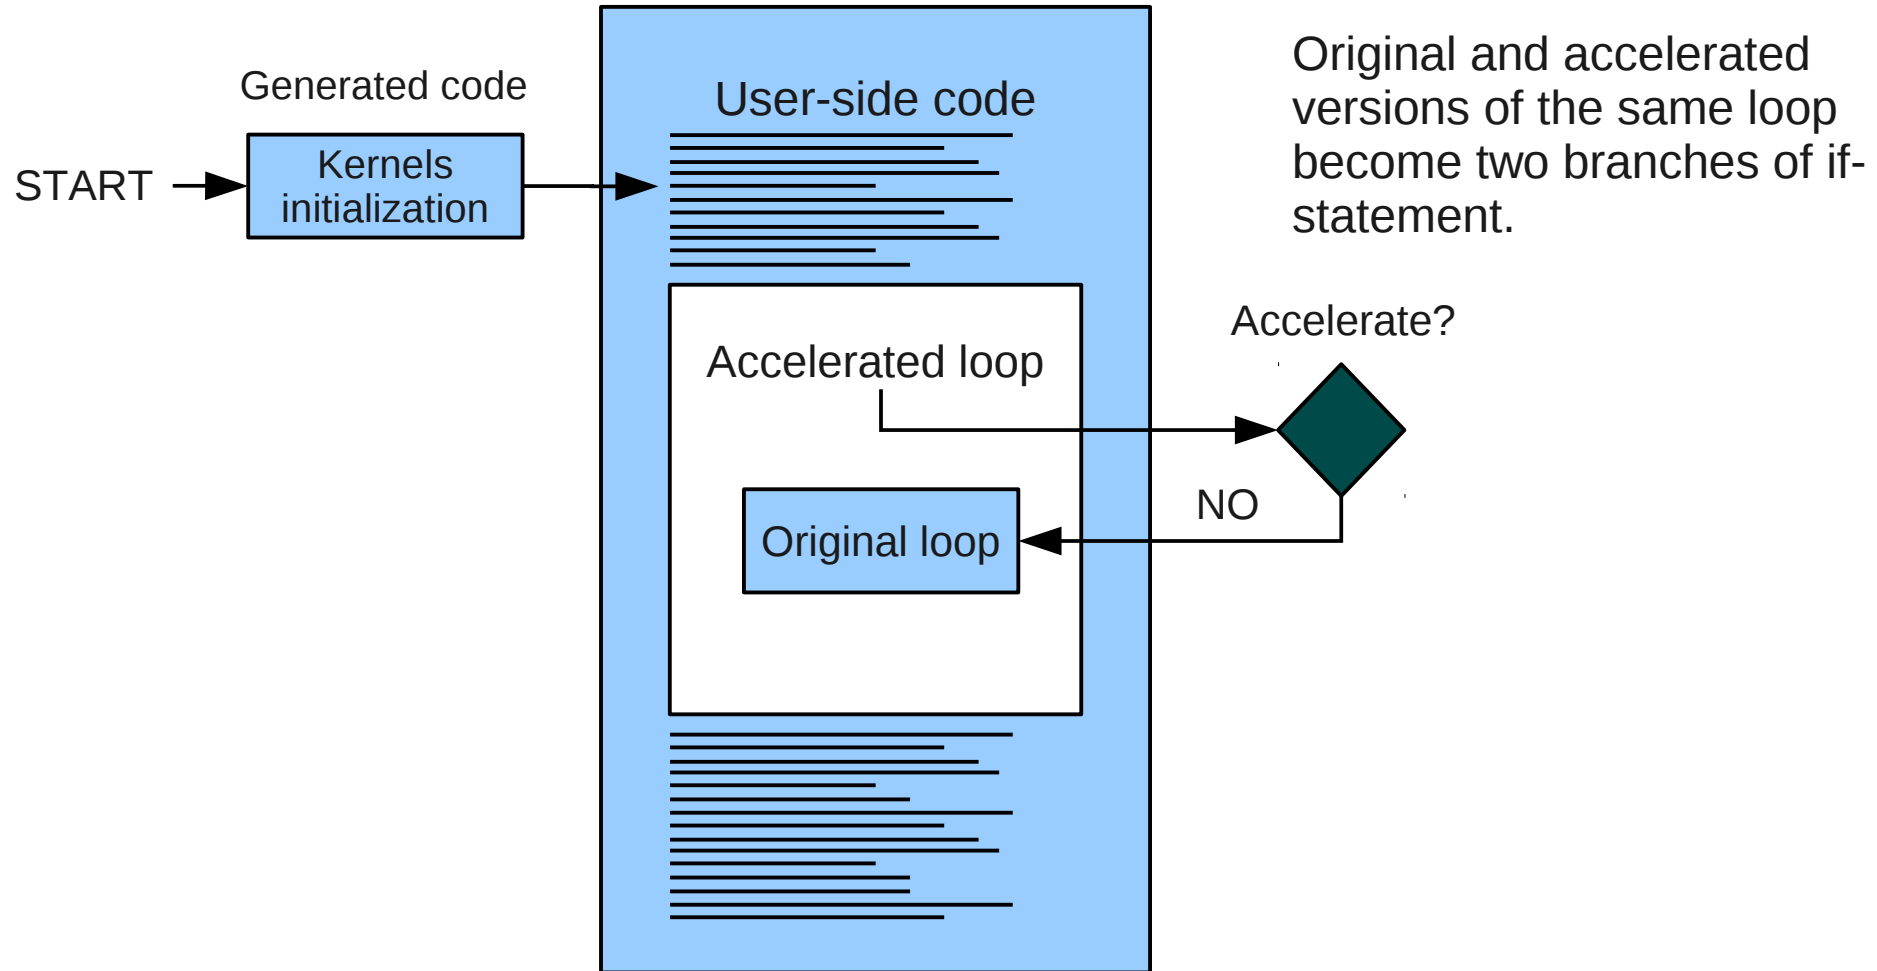


Equivalent device code is generated for suitable loops.

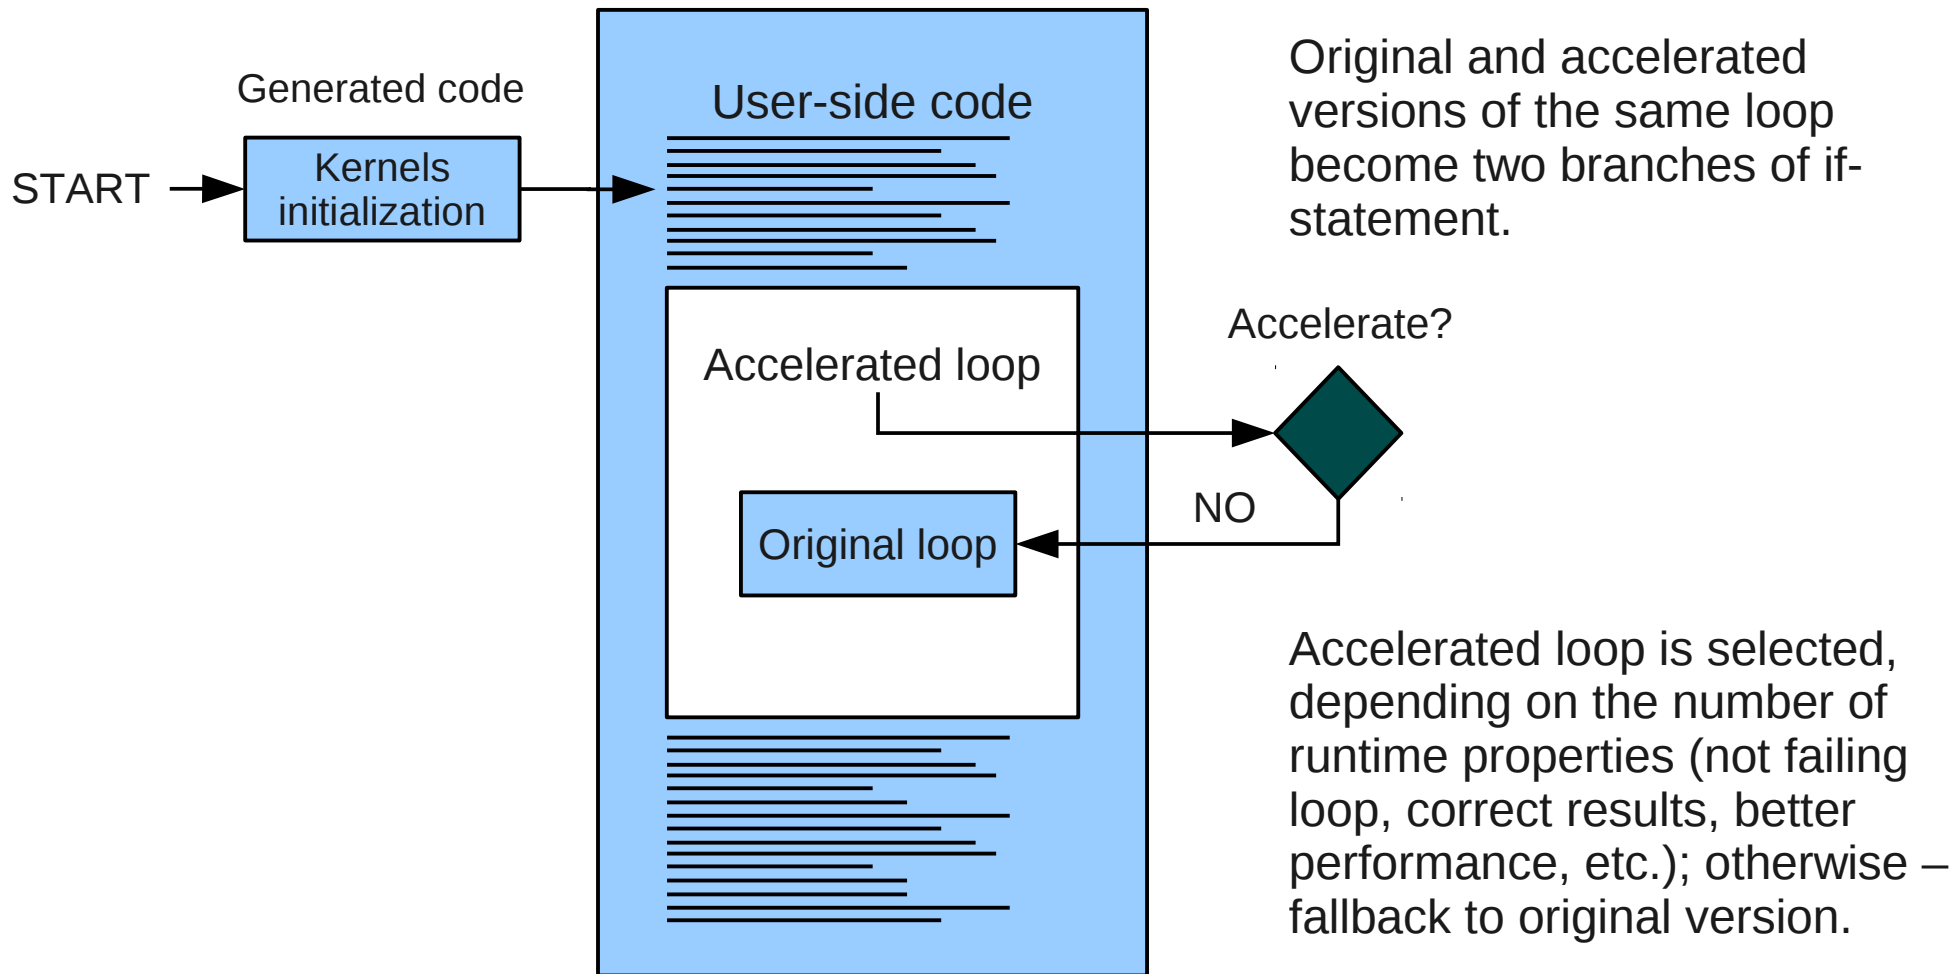
Additionally global constructors are generated to initialize configuration structures (with status, profiling, permanent dependencies, etc.) for each kernel.



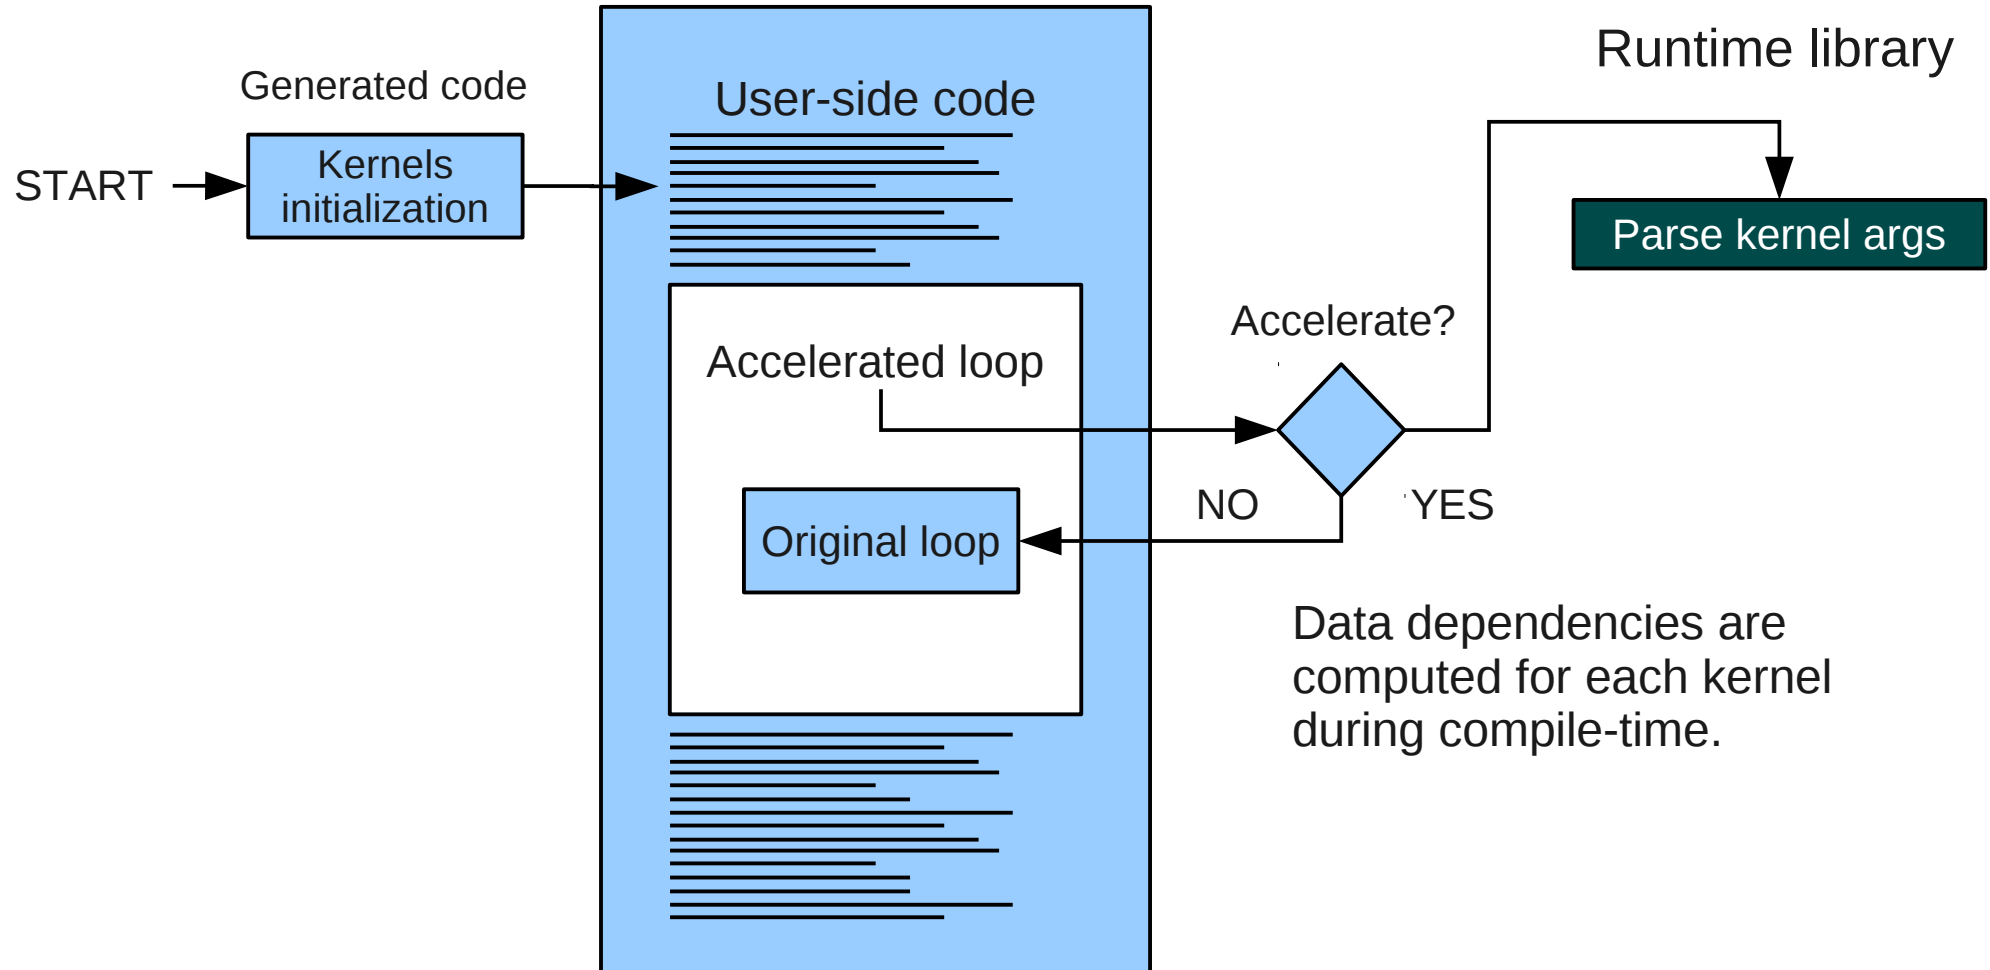
# Runtime workflow



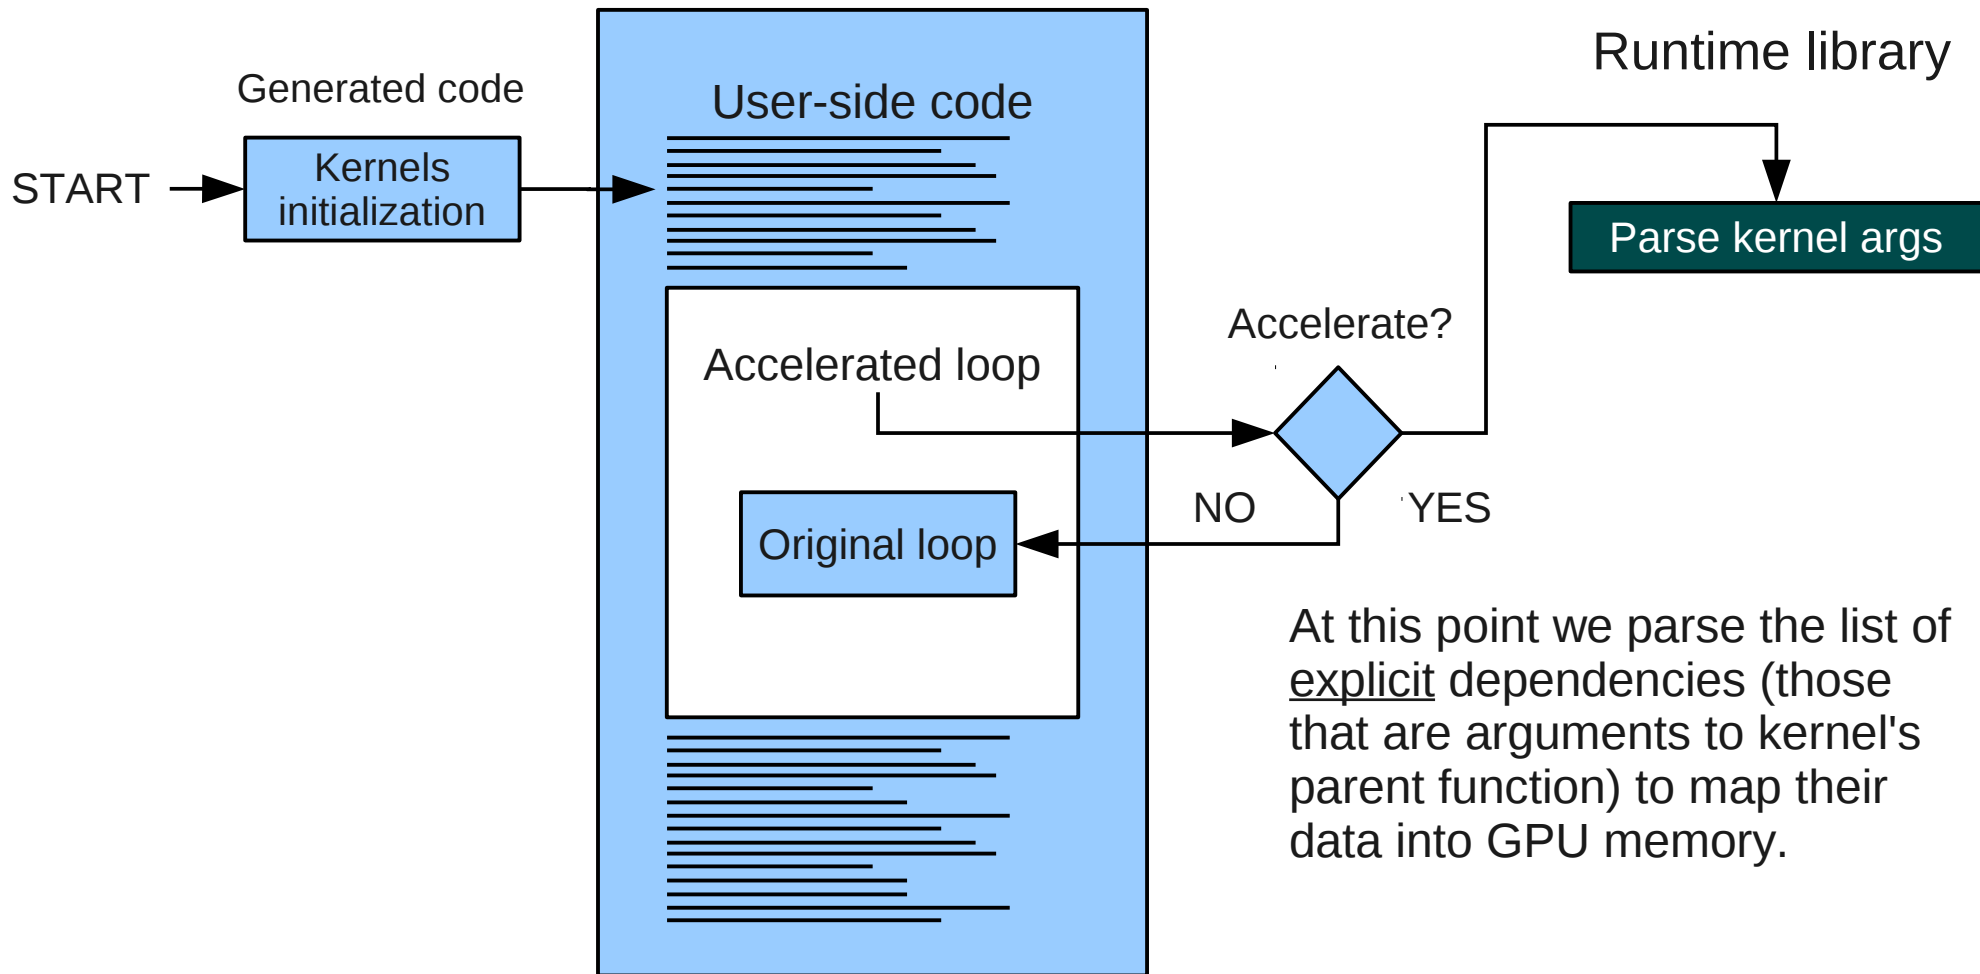
# Runtime workflow



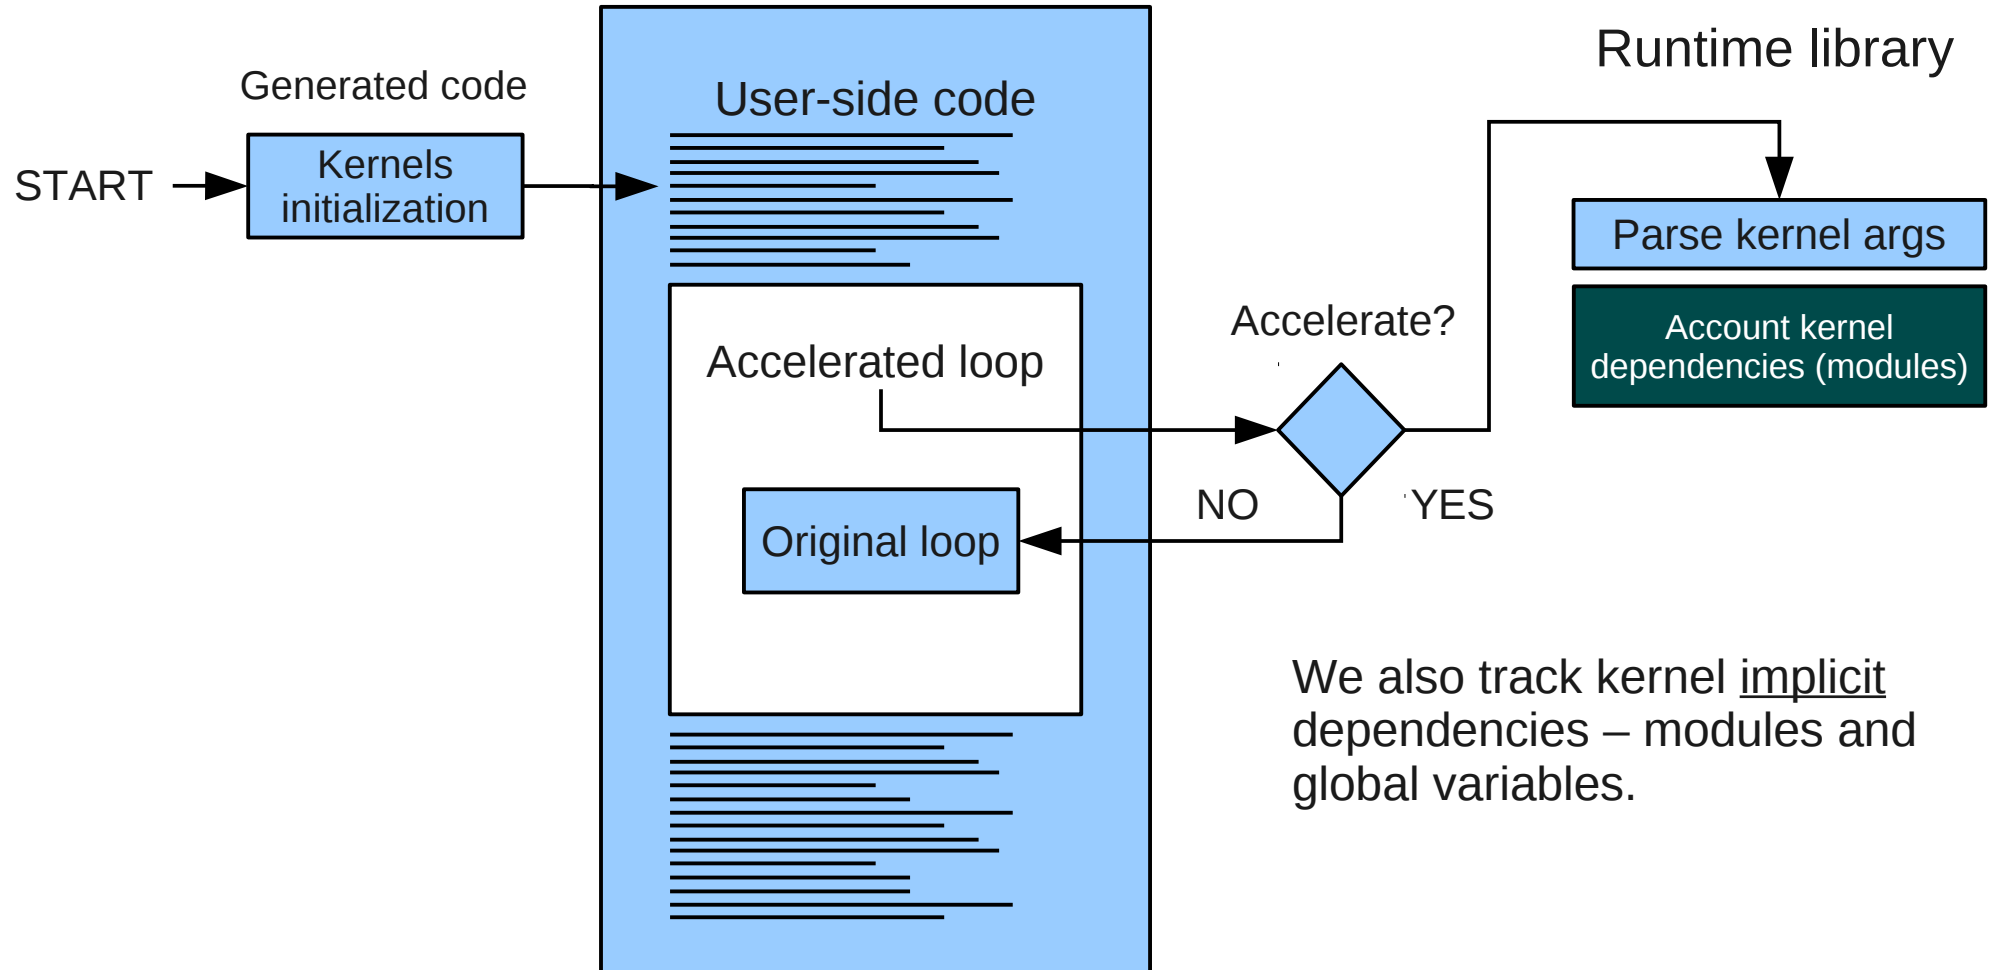
# Runtime workflow



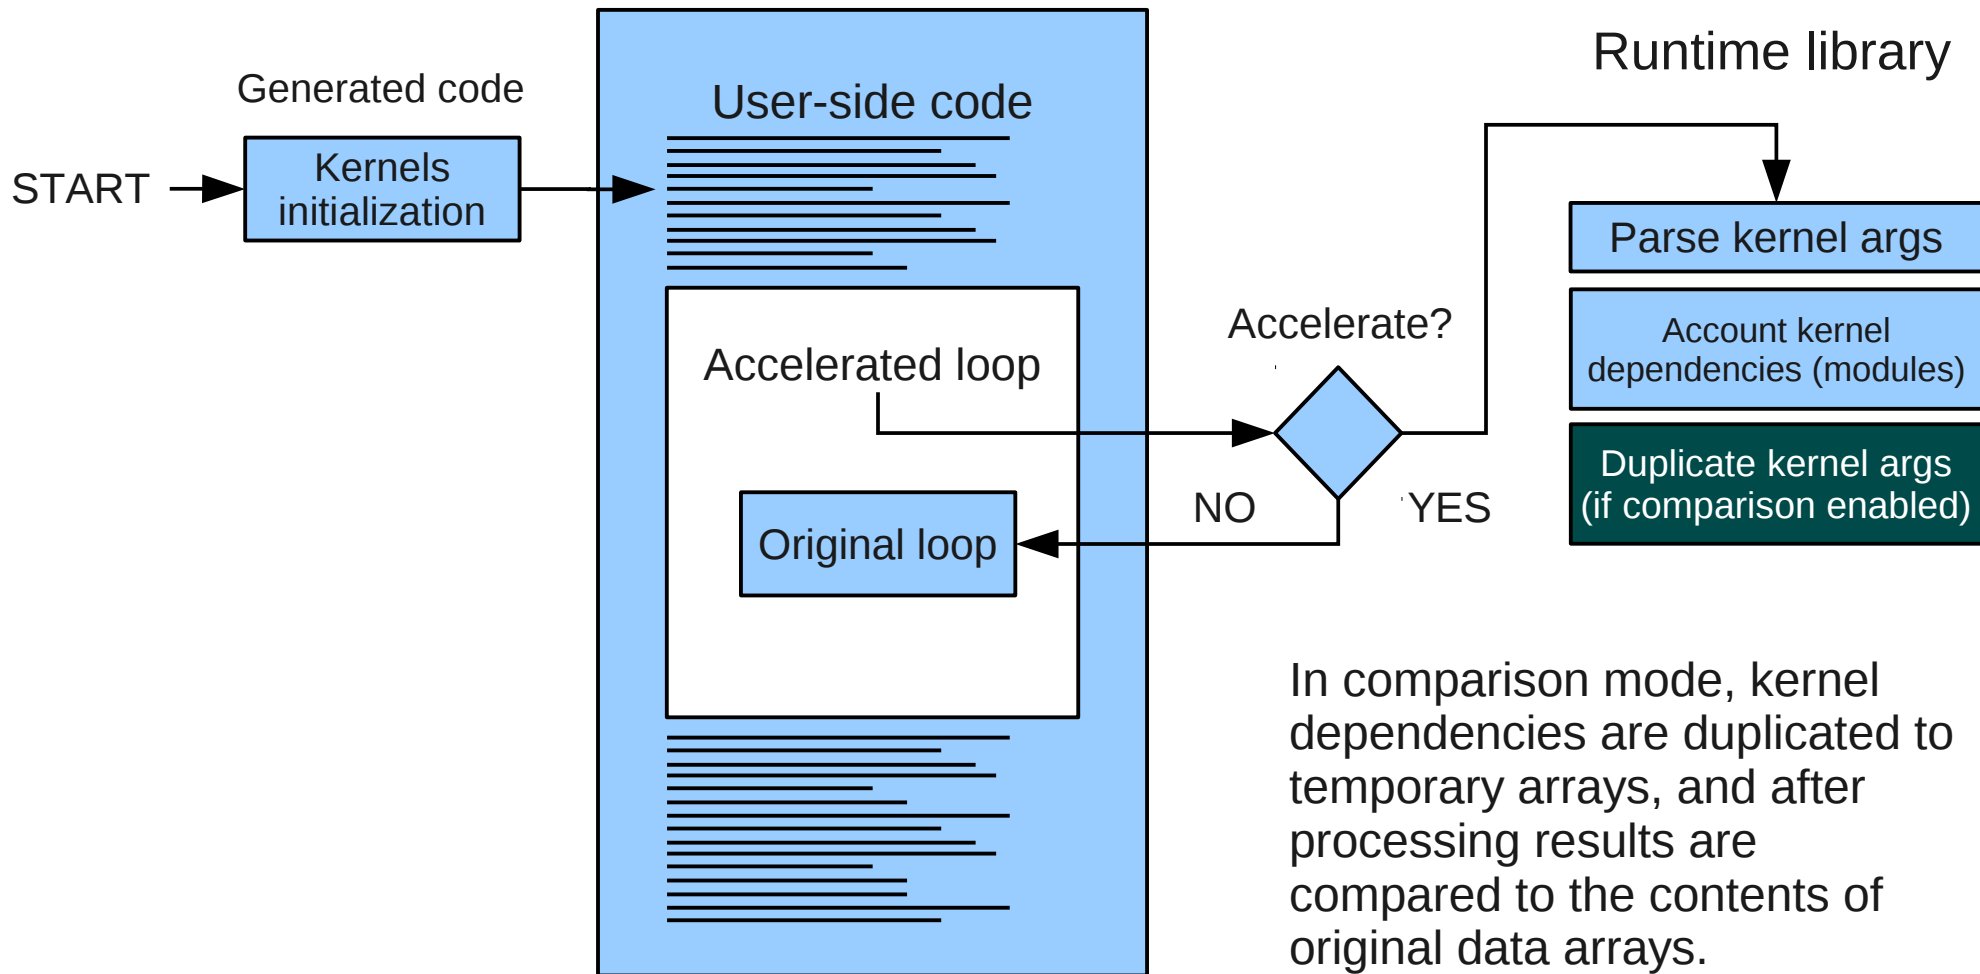
# Runtime workflow



# Runtime workflow

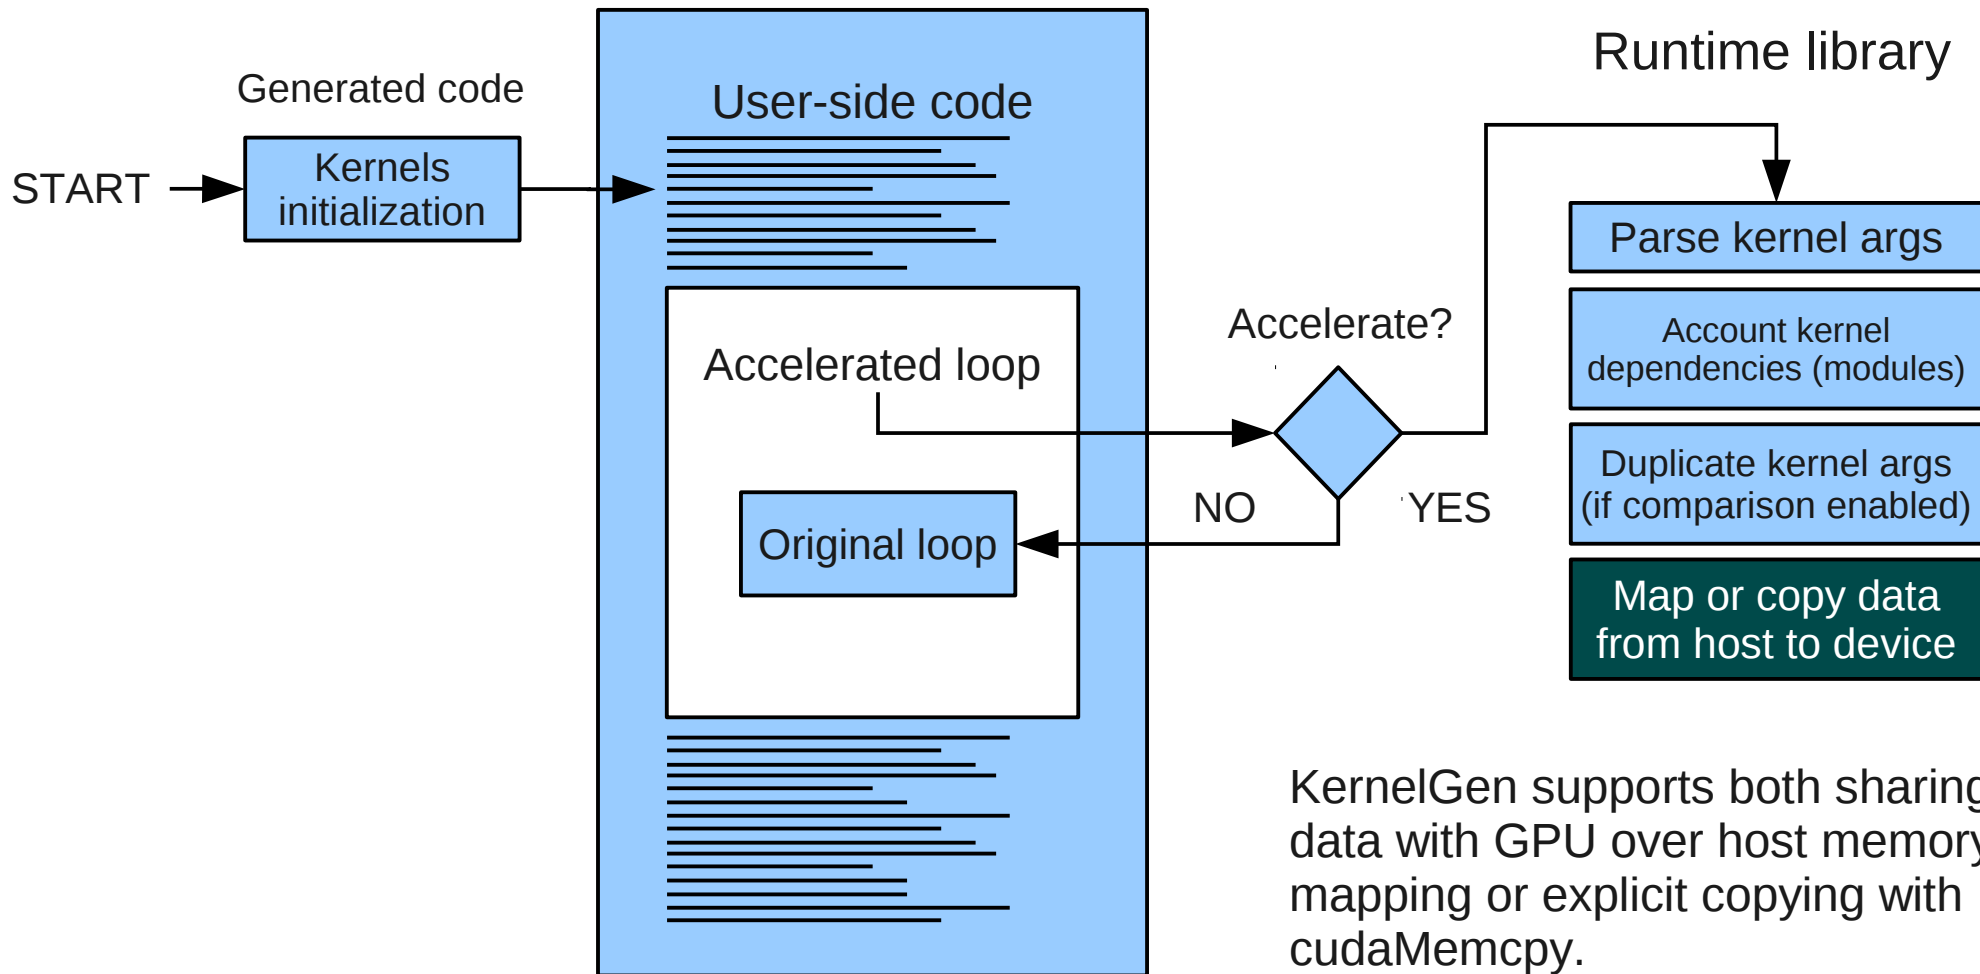


# Runtime workflow

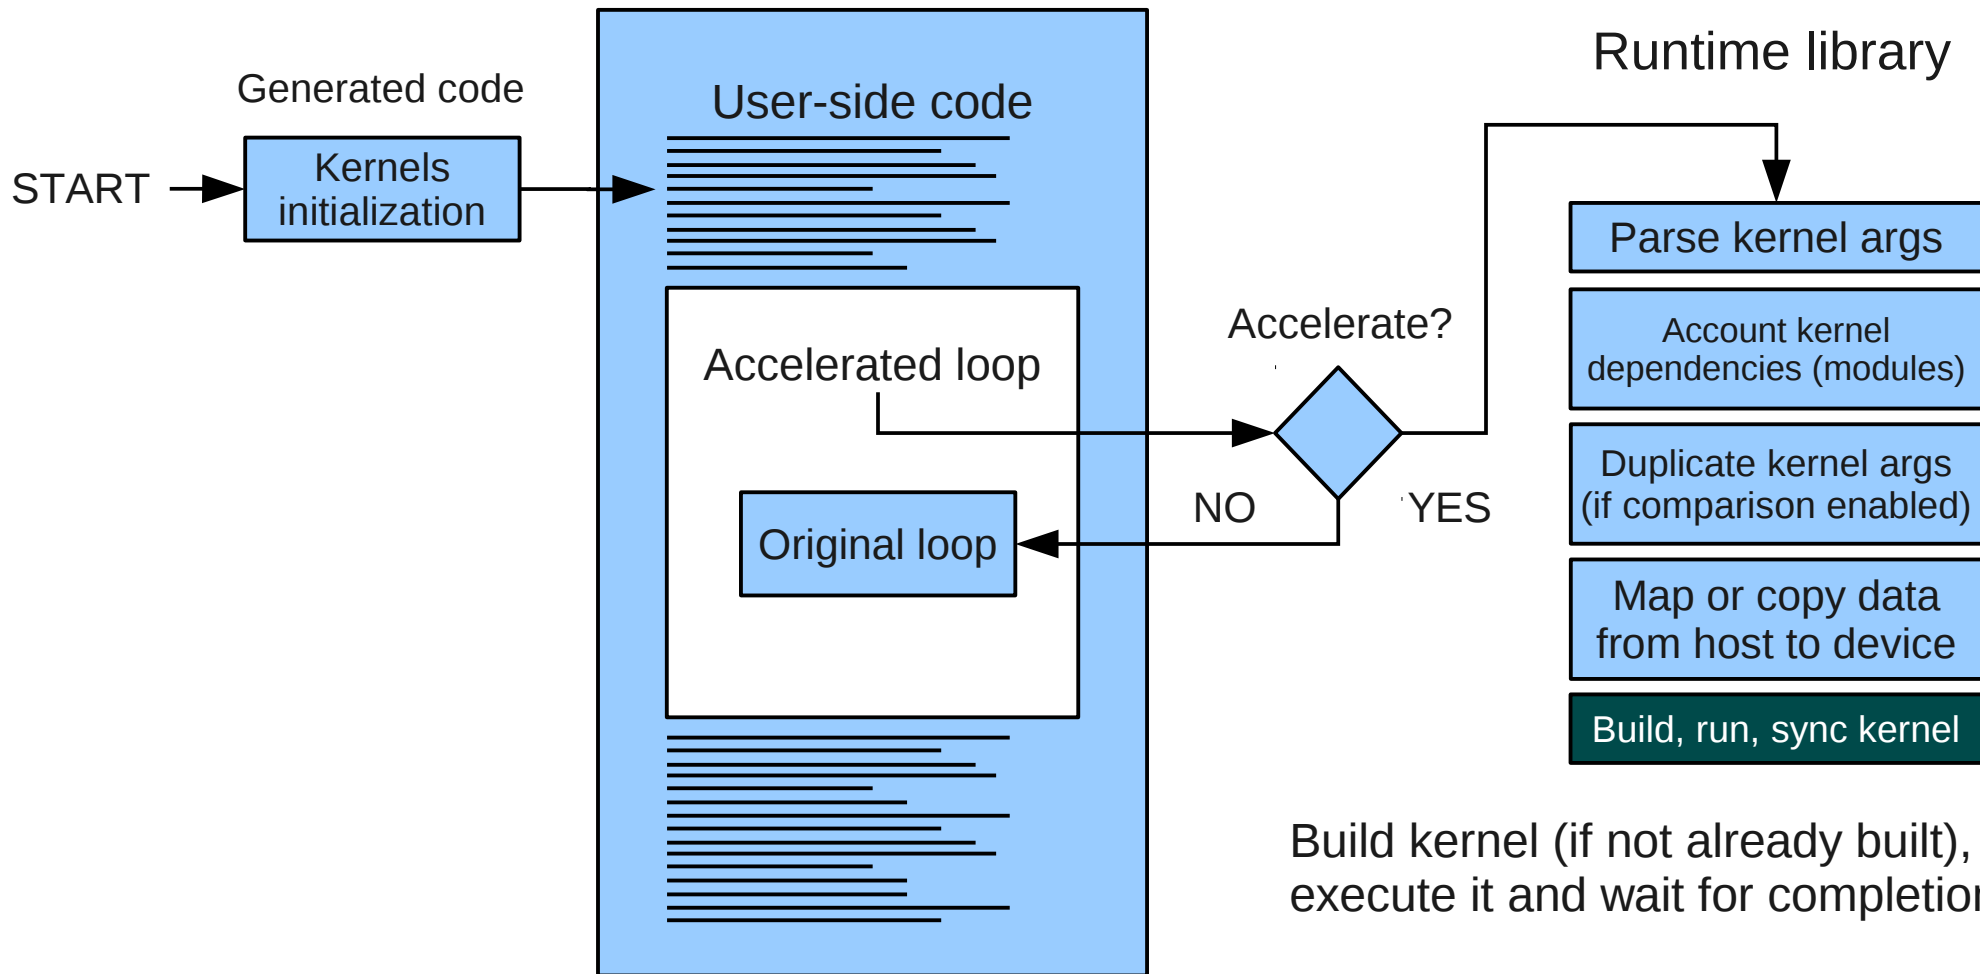




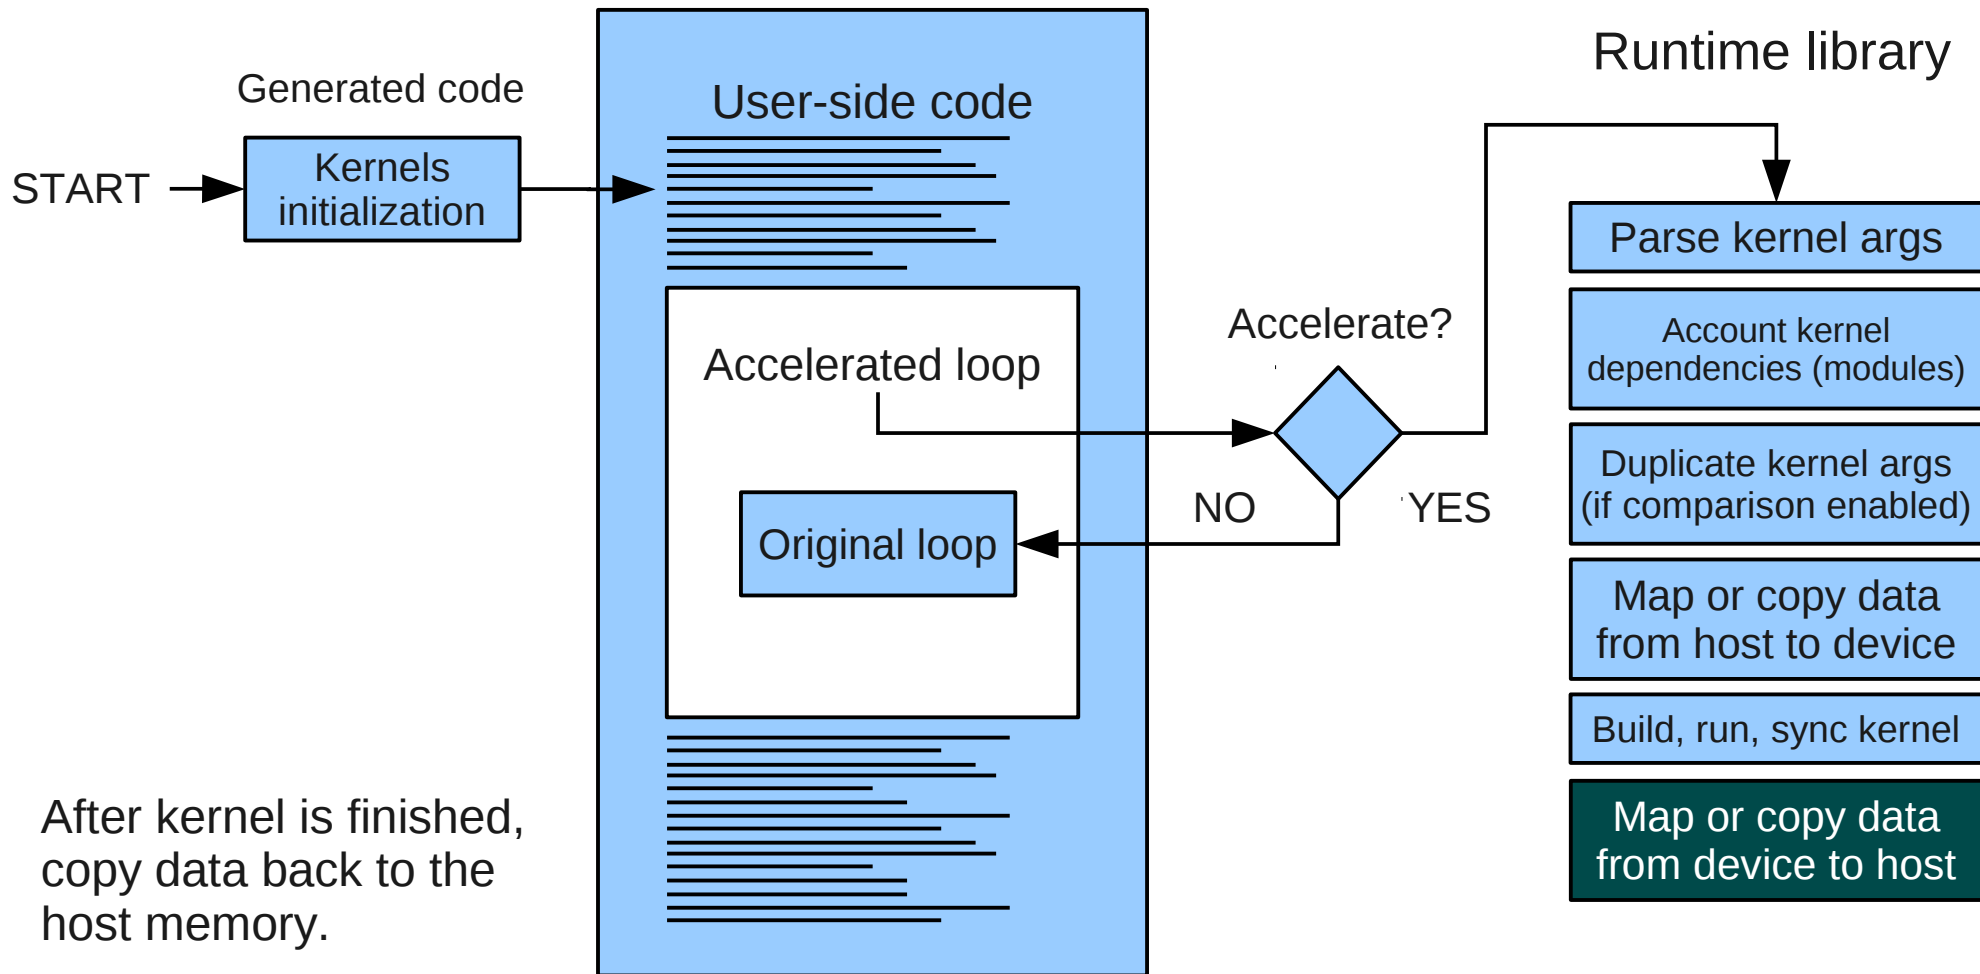
# Runtime workflow



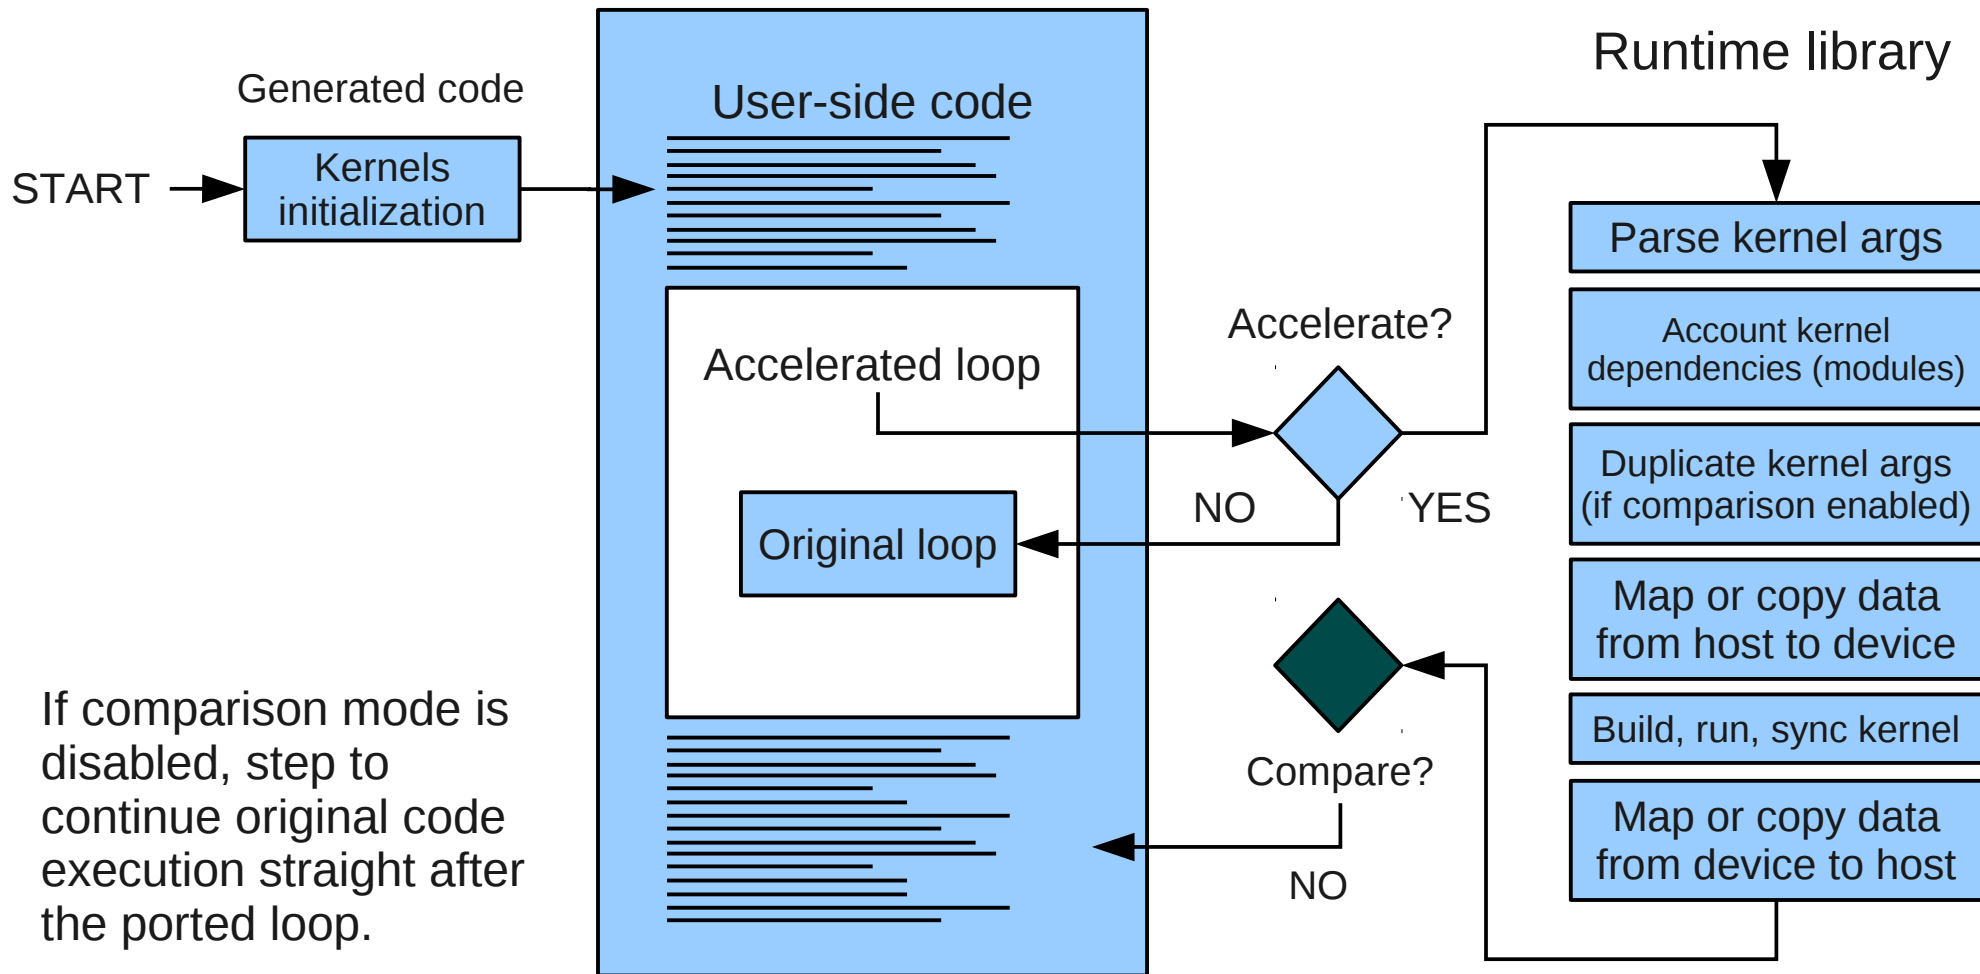
# Runtime workflow



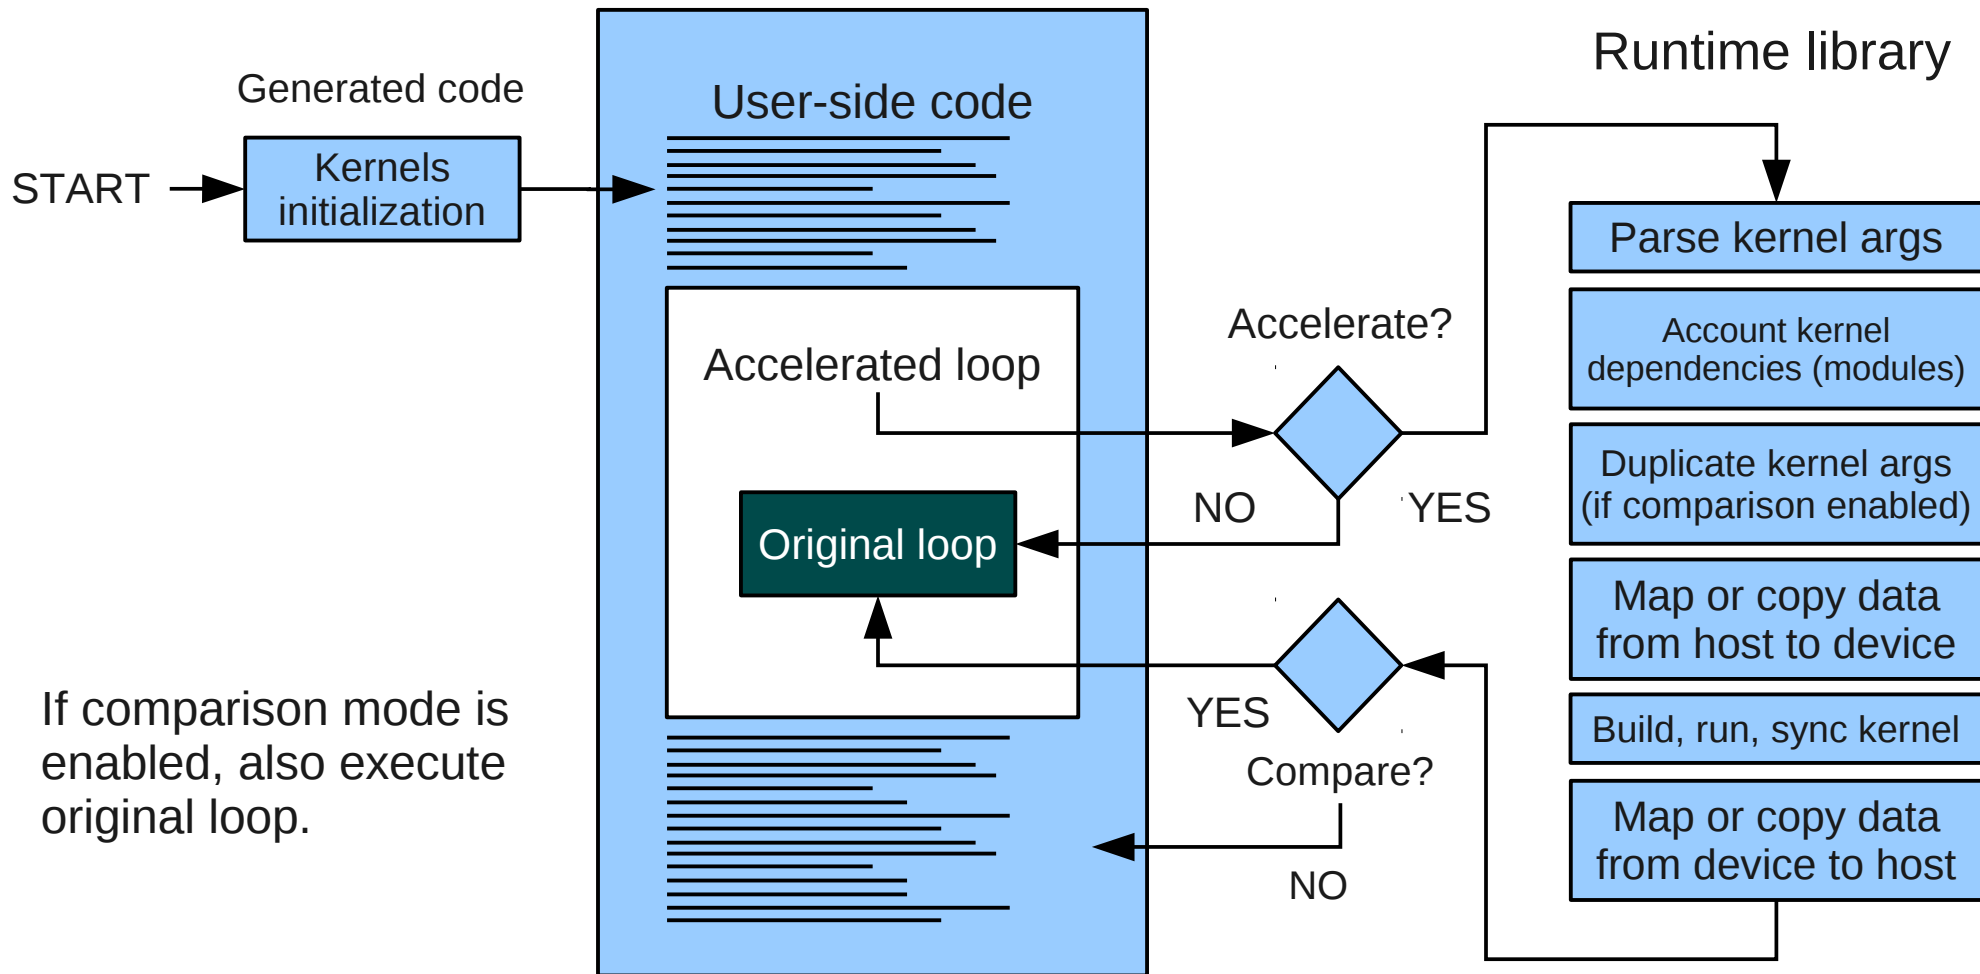
# Runtime workflow



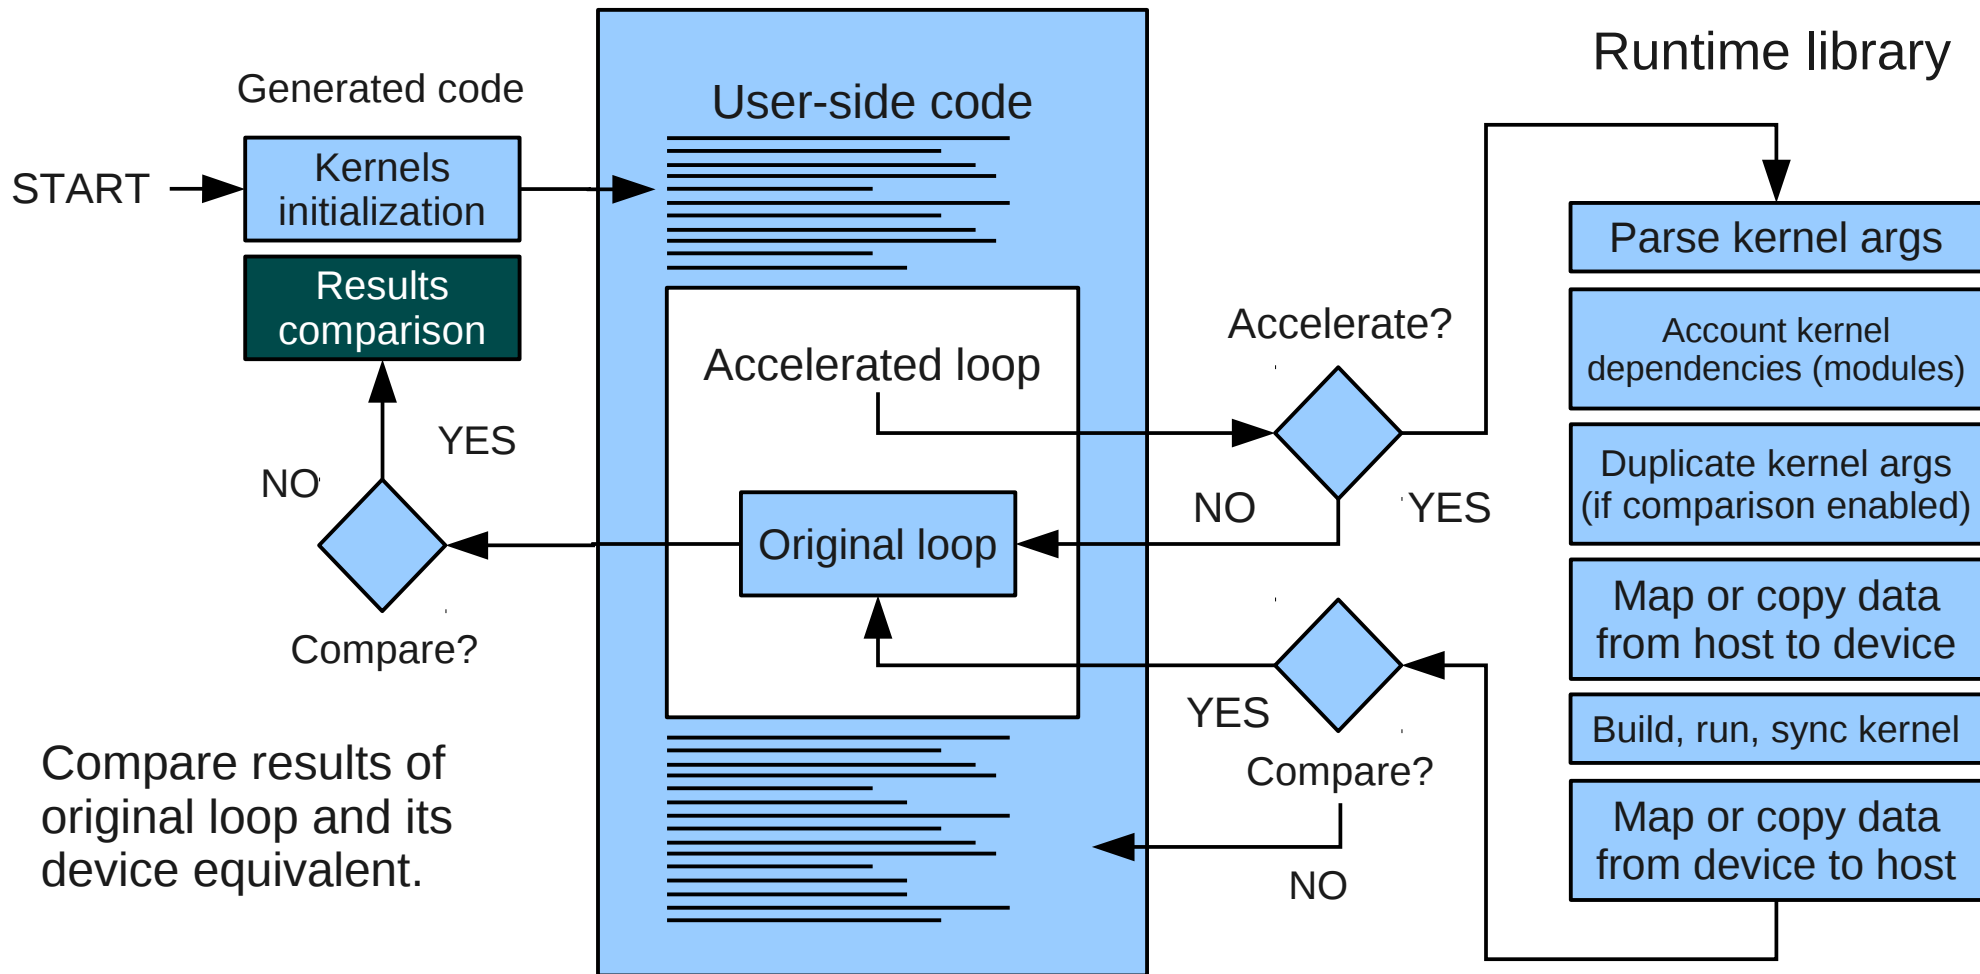
# Runtime workflow



# Runtime workflow

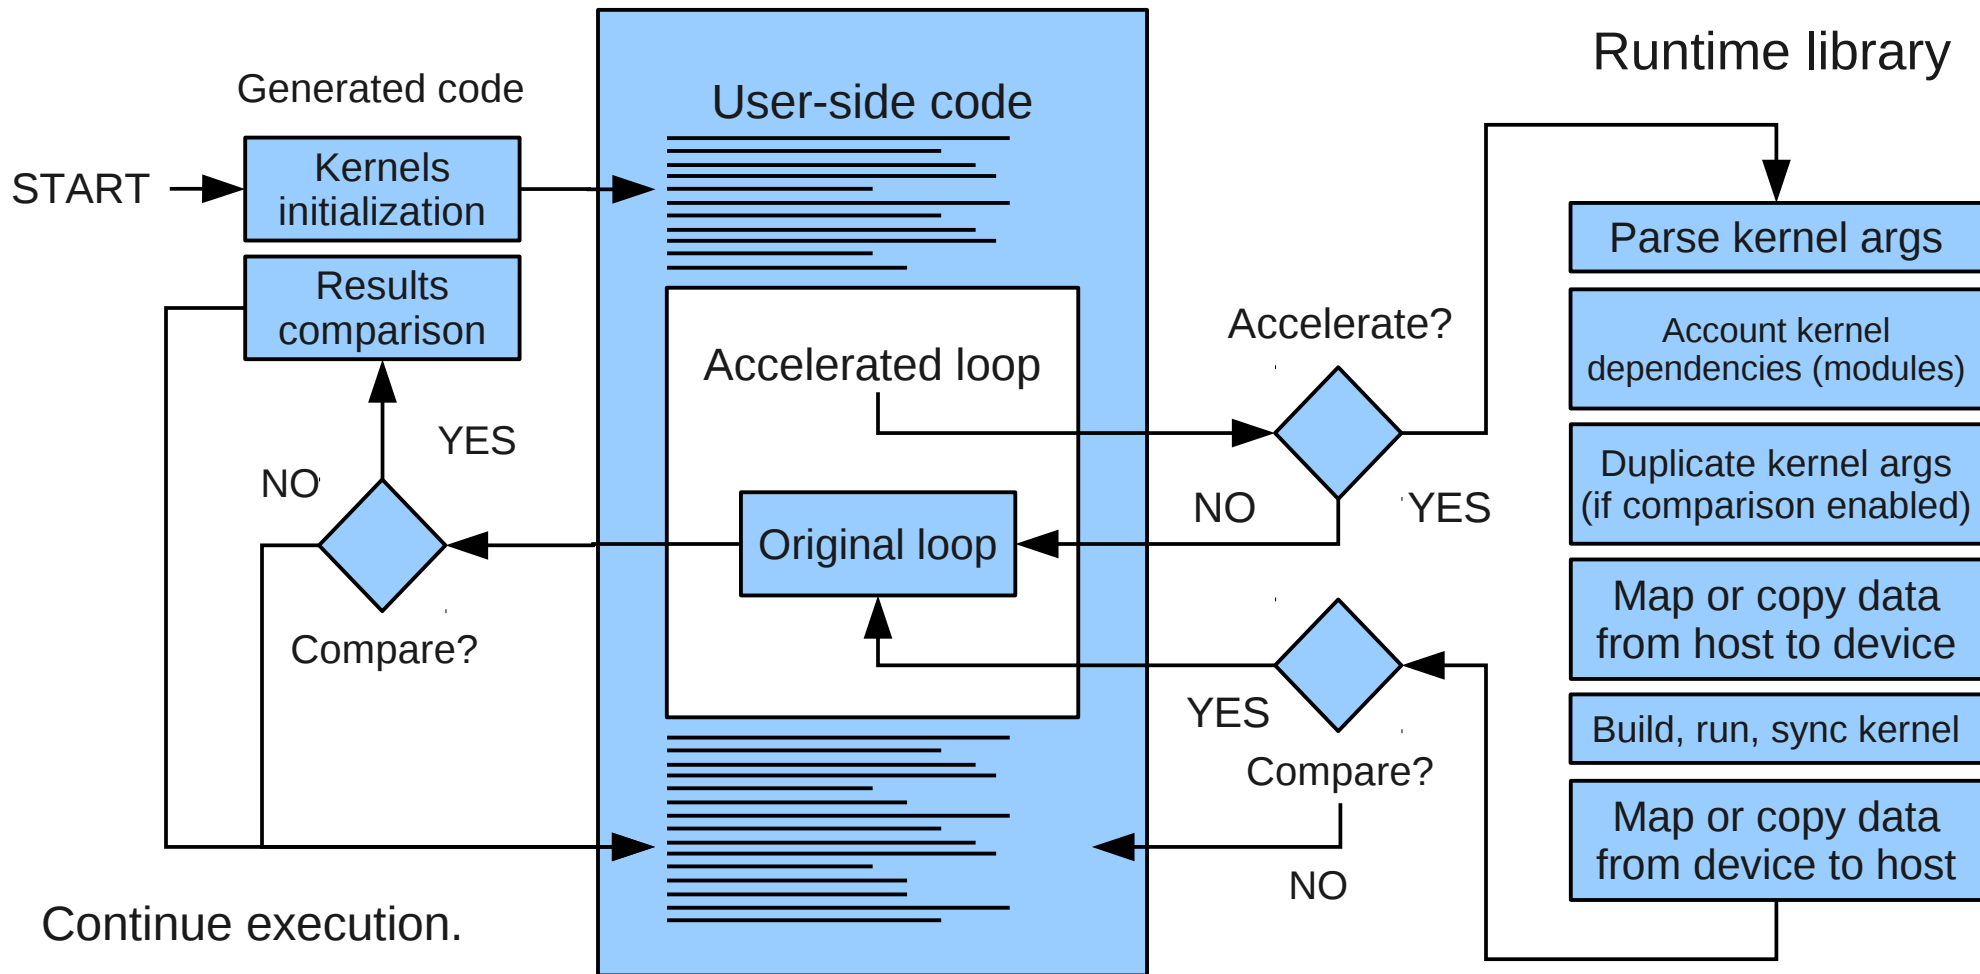


# Runtime workflow





# Runtime workflow

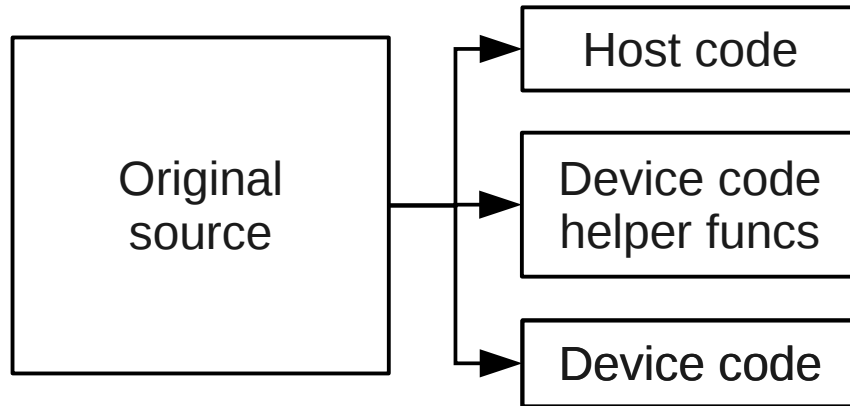


# Code generation workflow

Two parts of code generation process:

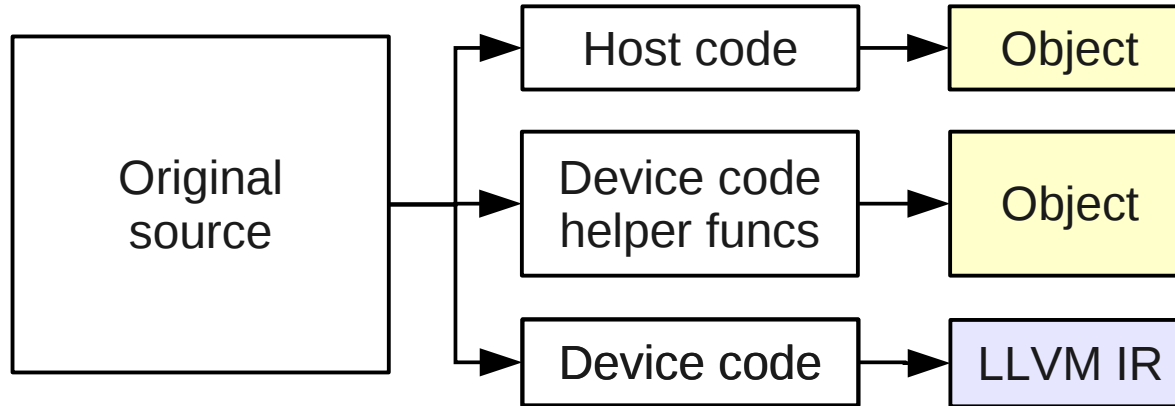
- **Compile time** – generate kernels strictly corresponding to original host loops
- **Runtime** – generate kernels, using additional info available at runtime: inline external functions, optimize compute grid, etc.

# Code generation workflow (compile-time part)



Loops suitable for device execution are identified in original source code, their bodies are surrounded with if-statement to switch between original loop and call to device kernel for this loop. Each suitable loop is duplicated in form of subroutine in a separate compilation unit. Additionally, helper initialization anchors are generated.

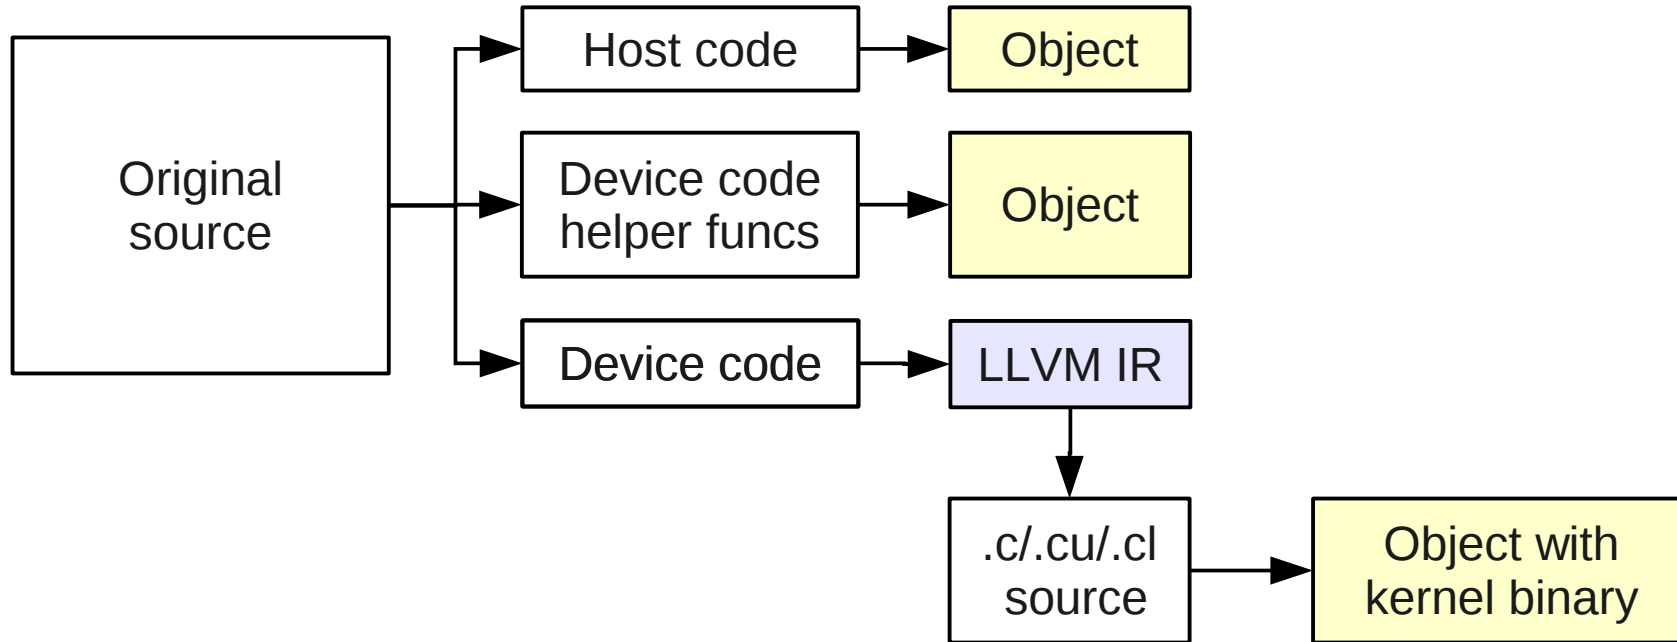
# Code generation workflow (compile-time part)



Objects for host code and device code helper functions can be generated directly with CPU compiler used by application.

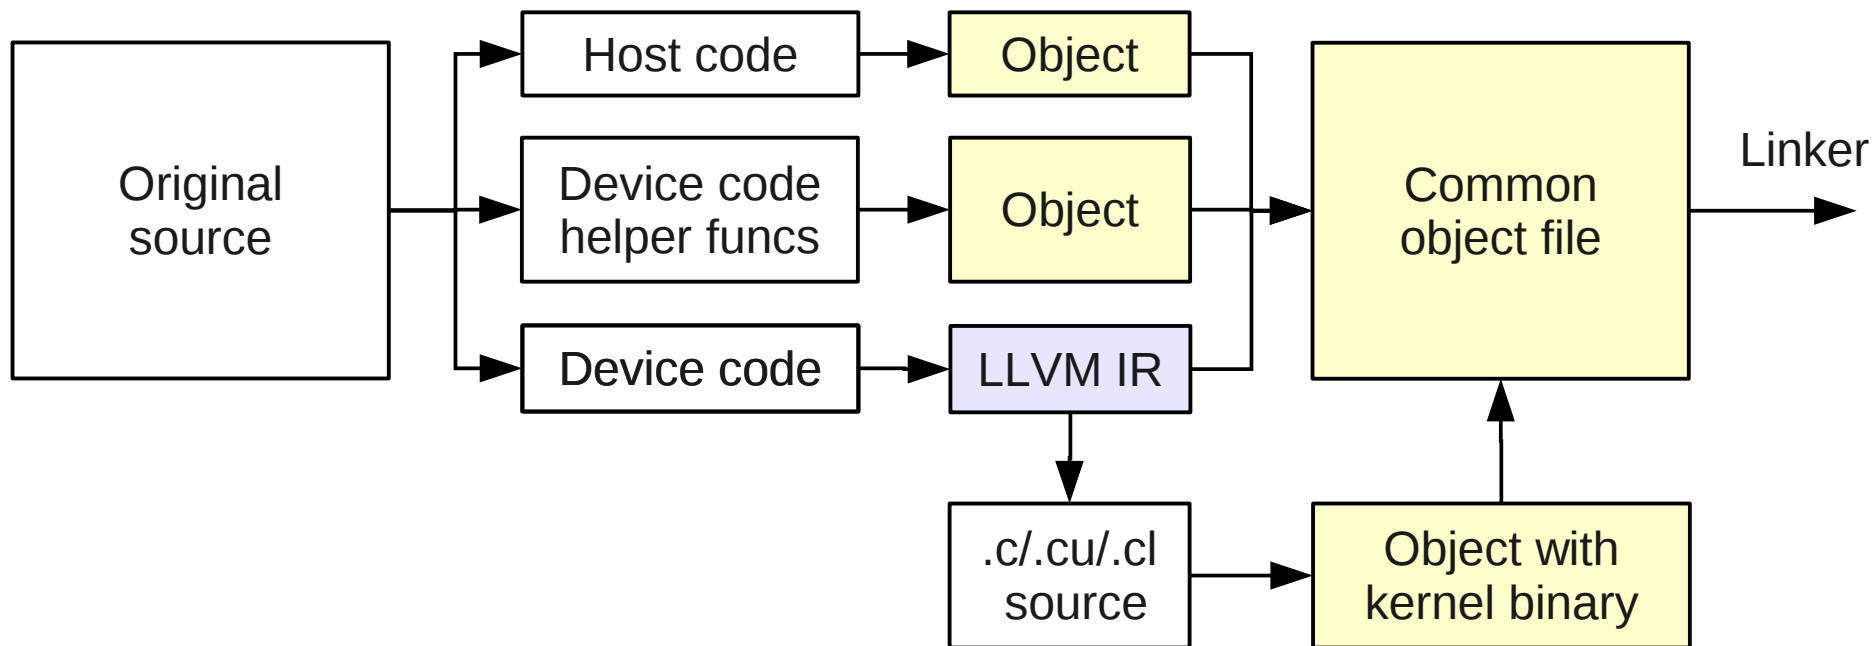
Device code is compiled into Low-Level Virtual Machine Intermediate representation (LLVM IR).

# Code generation workflow (compile-time part)



Code from LLVM IR is translated into C, CUDA or OpenCL using modified LLVM C Backend and compiled using the corresponding device compiler.

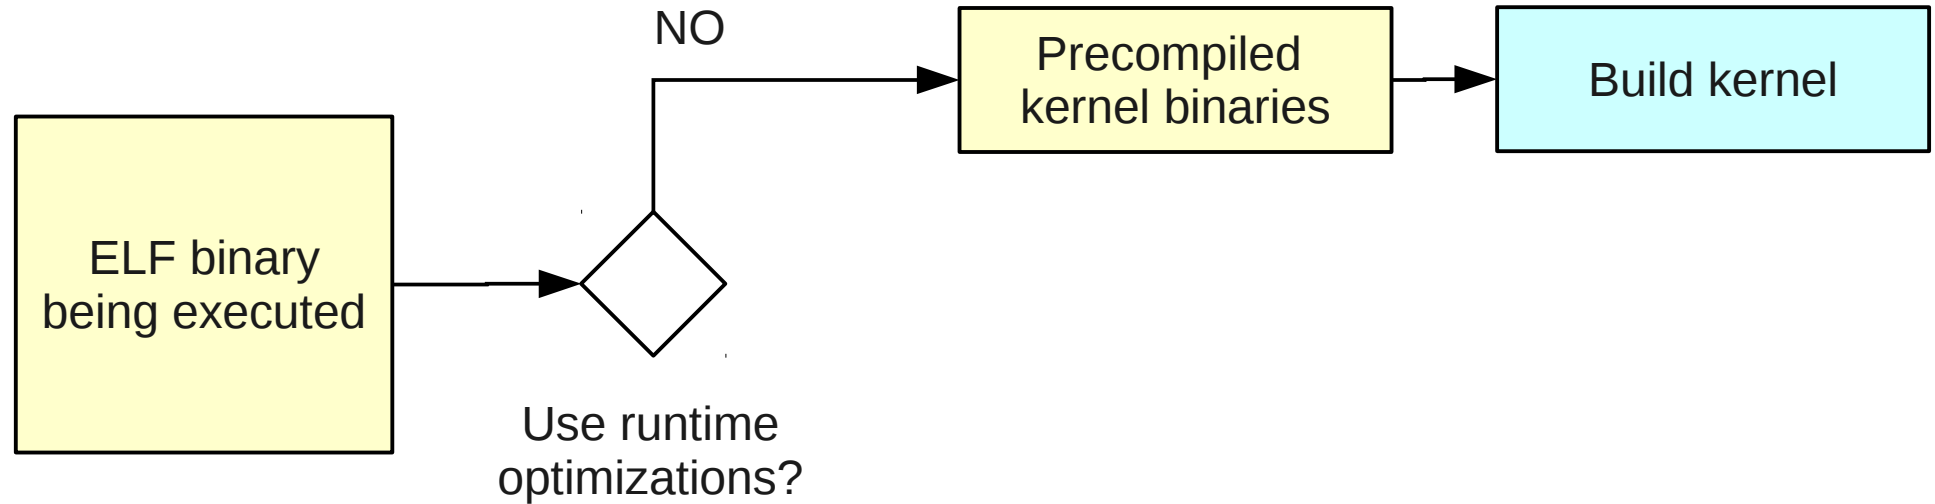
# Code generation workflow (compile-time part)



Finally, objects for all parts of the code are merged into single object to conserve “1 source → 1 object” layout. LLVM IR is also embedded into resulting object.

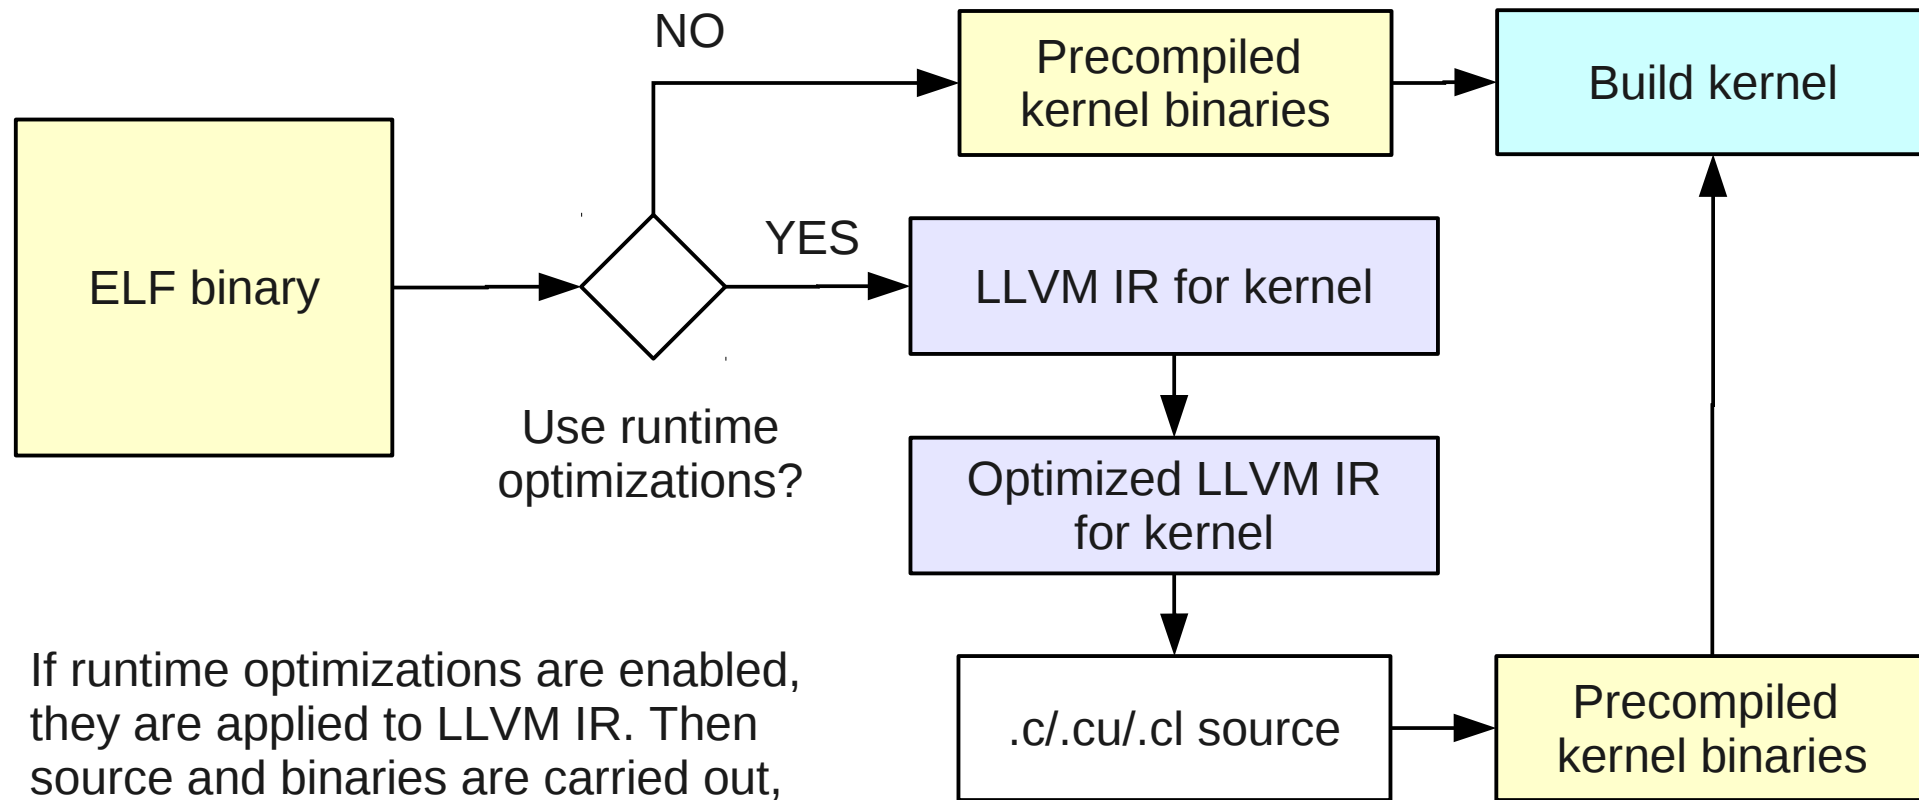


# Code generation workflow (runtime part)



Without runtime optimizations enabled, the previously compiled kernel binary could be built and executed.

# Code generation workflow (runtime part)



If runtime optimizations are enabled, they are applied to LLVM IR. Then source and binaries are carried out, just like in compile-time process.

## **3. Toolchain internals**

# Example: sincos

Consider toolchain steps in detail for the following simple test program:

```
subroutine sincos(nx, ny, nz, x, y, xy)

implicit none

integer, intent(in) :: nx, ny, nz
real, intent(in) :: x(nx, ny, nz), y(nx, ny, nz)
real, intent(inout) :: xy(nx, ny, nz)

integer :: i, j, k

do k = 1, nz
  do j = 1, ny
    do i = 1, nx
      xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
    enddo
  enddo
enddo

end subroutine sincos
```

# 1: host part of code split (1/3)

```
module sincos_kernelgen_module_uses
end module sincos_kernelgen_module_uses
module sincos_kernelgen_module
USE KERNELGEN

type(kernelgen_kernel_config), bind(C) :: sincos_loop_1_kernelgen_config

interface
function sincos_loop_1_kernelgen_compare()
end function

end interface

end module sincos_kernelgen_module

subroutine sincos(nx, ny, nz, x, y, xy)

USE KERNELGEN
USE sincos_kernelgen_module

implicit none
```

# 1: host part of code split (1/3)

```
module sincos_kernelgen_module_uses
end module sincos_kernelgen_module_uses
module sincos_kernelgen_module
USE KERNELGEN
```

```
type(kernelgen_kernel_config), bind(C) :: sincos_loop_1_kernelgen_config
```

```
interface
function sincos_loop_1_kernelgen_compare()
end function
```

Per-kernel config structure

```
end interface
```

```
end module sincos_kernelgen_module
```

```
subroutine sincos(nx, ny, nz, x, y, xy)
```

```
USE KERNELGEN
USE sincos_kernelgen_module
```

```
implicit none
```

# 1: host part of code split (1/3)

```
module sincos_kernelgen_module_uses
end module sincos_kernelgen_module_uses
module sincos_kernelgen_module
USE KERNELGEN
```

```
type(kernelgen_kernel_config), bind(C) :: sincos_loop_1_kernelgen_config
```

```
interface
function sincos_loop_1_kernelgen_compare()
end function
```

```
end interface
```

```
end module sincos_kernelgen_module
```

```
subroutine sincos(nx, ny, nz, x, y, xy)
```

```
USE KERNELGEN
USE sincos_kernelgen_module
```

```
implicit none
```

Adding kernel-specific and internal module with runtime calls

# 1: host part of code split (2/3)

```
!$KERNELGEN SELECT sincos_loop_1_kernelgen
if (sincos_loop_1_kernelgen_config%runmode .ne. kernelgen_runmode_host) then
!$KERNELGEN CALL sincos_loop_1_kernelgen
    call kernelgen_launch(sincos_loop_1_kernelgen_config, 1, nx, 1, ny, 1, nz, 6, 0,
nz, sizeof(nz), nz, ny, sizeof(ny), ny, nx, sizeof(nx), nx, xy, sizeof(xy), xy, x,
sizeof(x), x, y, sizeof(y), y)
    k = nz + 1
    j = ny + 1
    i = nx + 1
!$KERNELGEN END CALL sincos_loop_1_kernelgen
endif
if ((iand(sincos_loop_1_kernelgen_config%runmode, kernelgen_runmode_host) .eq. 1)
.or. (kernelgen_get_last_error() .ne. 0)) then
!$KERNELGEN LOOP sincos_loop_1_kernelgen
do k = 1, nz
    do j = 1, ny
        do i = 1, nx
            xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
        enddo
    enddo
enddo
!$KERNELGEN END LOOP sincos_loop_1_kernelgen
endif
```



# 1: host part of code split (2/3)

```
!$KERNELGEN SELECT sincos_loop_1 kernelgen
```

```
if (sincos_loop_1
```

```
!$KERNELGEN CALL
```

```
call kernelgen_
```

```
nz, sizeof(nz), nz, ny, sizeof(ny), ny, nx, sizeof(nx), nx, xy, sizeof(xy), xy, x,  
sizeof(x), x, y, sizeof(y), y)
```

```
k = nz + 1
```

```
j = ny + 1
```

```
i = nx + 1
```

```
!$KERNELGEN END CALL sincos_loop_1_kernelgen
```

```
endif
```

```
if ((iand(sincos_loop_1_kernelgen_config%runmode, kernelgen_runmode_host) .eq. 1)  
.or. (kernelgen_get_last_error() .ne. 0)) then
```

```
!$KERNELGEN LOOP sincos_loop_1_kernelgen
```

```
do k = 1, nz
```

```
do j = 1, ny
```

```
do i = 1, nx
```

```
xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
```

```
enddo
```

```
enddo
```

```
enddo
```

```
!$KERNELGEN END LOOP sincos_loop_1_kernelgen
```

```
endif
```

Loop location marker for processing script to clear everything here, if kernel was not successfully compiled.

0,

# 1: host part of code split (2/3)

```
!$KERNELGEN SELECT sincos_loop_1_kernelgen
if (sincos_loop_1_kernelgen_config%runmode .ne. kernelgen_runmode_host) then
!$KERNELGEN CALL sincos_loop_1_kernelgen
  call kernelgen_launch(
    nz, sizeof(nz), nz, ny, sizeof(ny), ny, nx, sizeof(nx), nx, xy, sizeof(xy), xy, x,
    sizeof(x), x, y, sizeof(y), y)
  k = nz + 1
  j = ny + 1
  i = nx + 1
!$KERNELGEN END CALL sincos_loop_1_kernelgen
endif
if ((iand(sincos_loop_1_kernelgen_config%runmode, kernelgen_runmode_host) .eq. 1)
.or. (kernelgen_get_last_error() .ne. 0)) then
!$KERNELGEN LOOP sincos_loop_1_kernelgen
do k = 1, nz
  do j = 1, ny
    do i = 1, nx
      xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
    enddo
  enddo
enddo
!$KERNELGEN END LOOP sincos_loop_1_kernelgen
endif
```

If kernel is requested to be executed not only on host

# 1: host part of code split (2/3)

```
!$KERNELGEN SELECT sincos_loop_1_kernelgen
if (sincos_loop_1_kernelgen_config%runmode .ne. kernelgen_runmode_host) then
!$KERNELGEN CALL sincos_loop_1_kernelgen
    call kernelgen_launch(sincos_loop_1_kernelgen_config, 1, nx, 1, ny, 1, nz, 6, 0,
nz, sizeof(nz), nz, ny, sizeof(ny), ny, nx, sizeof(nx), nx, xy, sizeof(xy), xy, x,
sizeof(x), x, y, sizeof(y), y)
    k = nz + 1
    j = ny + 1
    i = nx + 1
!$KERNELGEN END CALL sincos_loop_1_kernelgen
endif
if ((iand(sincos_loop_1_kernelgen_config%runmode, kernelgen_runmode_host) .eq. 1)
.or. (kernelgen_get_last_error() .ne. 0)) then
!$KERNELGEN LOOP sincos_loop_1_kernelgen
do k = 1, nz
    do j = 1, ny
        do i = 1, nx
            xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
        enddo
    enddo
enddo
!$KERNELGEN END LOOP sincos_loop_1_kernelgen
endif
```

Launch kernel with its config handle, grid and dependencies

# 1: host part of code split (2/3)

```
!$KERNELGEN SELECT sincos_loop_1_kernelgen
if (sincos_loop_1_kernelgen_config%runmode .ne. kernelgen_runmode_host) then
!$KERNELGEN CALL sincos_loop_1_kernelgen
  call kernelgen_launch(sincos_loop_1_kernelgen_config, 1, nx, 1, ny, 1, nz, 6, 0,
    nz, sizeof(nz), nz, ny, sizeof(ny), ny, nx, sizeof(nx), nx, xy, sizeof(xy), xy, x,
    sizeof(x), x, y, sizeof(y), y)
  k = nz + 1
  j = ny + 1
  i = nx + 1
!$KERNELGEN
endif
if ((iand(sincos_loop_1_kernelgen_config%runmode, kernelgen_runmode_host) .eq. 1)
.or. (kernelgen_get_last_error() .ne. 0)) then
!$KERNELGEN LOOP sincos_loop_1_kernelgen
do k = 1, nz
  do j = 1, ny
    do i = 1, nx
      xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
    enddo
  enddo
enddo
!$KERNELGEN END LOOP sincos_loop_1_kernelgen
endif
```

Just in case increment old indexes, like if they were used by loop

# 1: host part of code split (2/3)

```
!$KERNELGEN SELECT sincos_loop_1_kernelgen
if (sincos_loop_1_kernelgen_config%runmode .ne. kernelgen_runmode_host) then
!$KERNELGEN CALL sincos_loop_1_kernelgen
    call kernelgen_launch(sincos_loop_1_kernelgen_config, 1, nx, 1, ny, 1, nz, 6, 0,
nz, sizeof(nz), nz, ny, sizeof(ny), ny, nx, sizeof(nx), nx, xy, sizeof(xy), xy, x,
sizeof(x), x, y, sizeof(y), y)
    k = nz + 1
    j = ny + 1
    i = nx + 1
!$KERNELGEN END CALL sincos_loop_1_kernelgen
endif
if ((iand(sincos_loop_1_kernelgen_config%runmode, kernelgen_runmode_host) .eq. 1)
.or. (kernelgen_get_last_error() .ne. 0)) then
!$KERNELGEN
do k = 1, nz
    do j = 1, ny
        do i = 1, nx
            xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
        enddo
    enddo
enddo
!$KERNELGEN END LOOP sincos_loop_1_kernelgen
endif
```

If kernel is requested to be executed not only on host  
or there is an error executing kernel on device

# 1: host part of code split (2/3)

```
!$KERNELGEN SELECT sincos_loop_1_kernelgen
if (sincos_loop_1_kernelgen_config%runmode .ne. kernelgen_runmode_host) then
!$KERNELGEN CALL sincos_loop_1_kernelgen
    call kernelgen_launch(sincos_loop_1_kernelgen_config, 1, nx, 1, ny, 1, nz, 6, 0,
nz, sizeof(nz), nz, ny, sizeof(ny), ny, nx, sizeof(nx), nx, xy, sizeof(xy), xy, x,
sizeof(x), x, y, sizeof(y), y)
k = nz + 1
j = ny + 1
i = nx + 1
!$KERNELGEN END CALL sincos_loop_1_kernelgen
endif
if ((and(sincos_loop_1_kernelgen_config%runmode, kernelgen_runmode_host) .eq. 1)
.or. (k .eq. 0)) then
!$KERNELGEN LOOP sincos_loop_1_kernelgen
    do k = 1, nz
        do j = 1, ny
            do i = 1, nx
                xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
            enddo
        enddo
    enddo
!$KERNELGEN END LOOP sincos_loop_1_kernelgen
endif
```

Execute original loop

# 1: host part of code split (3/3)

```
if ((sincos_loop_1_kernelgen_config%compare .eq. 1) .and. (kernelgen_get_last_error()  
.eq. 0)) then  
  call kernelgen_compare(sincos_loop_1_kernelgen_config,  
sincos_loop_1_kernelgen_compare, kernelgen_compare_maxdiff)
```

```
endif
```

```
!$KERN
```

If no error and comparison enabled, compare results of CPU and device

```
end subroutine sincos
```

## 2: device part of code split (1/2)

```
subroutine sincos_loop_1_kernelgen(nz, ny, nx, xy, x, y)
```

```
implicit none
```

```
interface
```

Kernel subroutine name is a decorated name of original loop function

```
subroutine sincos_loop_1_kernelgen_blockidx_x(index, start, end) bind(C)
```

```
use iso_c_binding
```

```
integer(c_int) :: index
```

```
integer(c_int), value :: start, end
```

```
end subroutine
```

```
subroutine sincos_loop_1_kernelgen_blockidx_y(index, start, end) bind(C)
```

```
use iso_c_binding
```

```
integer(c_int) :: index
```

```
integer(c_int), value :: start, end
```

```
end subroutine
```

```
subroutine sincos_loop_1_kernelgen_blockidx_z(index, start, end) bind(C)
```

```
use iso_c_binding
```

```
integer(c_int) :: index
```

```
integer(c_int), value :: start, end
```

```
end subroutine
```

```
end interface
```



## 2: device part of code split (1/2)

```
subroutine sincos_loop_1_kernelgen(nz, ny, nx, xy, x, y)
implicit none
interface
  subroutine sincos_loop_1_kernelgen_blockidx_x(index, start, end) bind(C)
  use iso_c_binding
  integer(c_int) :: index
  integer(c_int), value :: start, end
  end subroutine
  subroutine sincos_loop_1_kernelgen_blockidx_y(index, start, end) bind(C)
  use iso_c_binding
  integer(c_int) :: index
  integer(c_int), value :: start, end
  end subroutine
  subroutine sincos_loop_1_kernelgen_blockidx_z(index, start, end) bind(C)
  use iso_c_binding
  integer(c_int) :: index
  integer(c_int), value :: start, end
  end subroutine
end interface
```

Interfaces to device functions returning device compute grid dimensions

## 2: device part of code split (2/2)

```
__device__ void sincos_loop_1_kernelgen_blockidx_z(k, 1, nz)
{
    #ifdef __CUDA_DEVICE_FUNC__
    call sincos_loop_1_kernelgen_blockidx_z(k, 1, nz)
    #else
    do k = 1, nz
    #endif
    #ifdef __CUDA_DEVICE_FUNC__
    call sincos_loop_1_kernelgen_blockidx_y(j, 1, ny)
    #else
    do j = 1, ny
    #endif
    #ifdef __CUDA_DEVICE_FUNC__
    call sincos_loop_1_kernelgen_blockidx_x(i, 1, nx)
    #else
    do i = 1, nx
    #endif
        xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
    #ifndef __CUDA_DEVICE_FUNC__
    enddo
    #endif
    #ifndef __CUDA_DEVICE_FUNC__
    enddo
    #endif
    #ifndef __CUDA_DEVICE_FUNC__
    enddo
    #endif
end subroutine sincos_loop_1_kernelgen
```

In device kernels loops indexes are computed using block/thread indexes

## 2: device part of code split (2/2)

```
#ifdef __CUDA_DEVICE_FUNC__
call sincos_loop_1_kernelgen_blockidx_z(k, 1, nz)
#else
do k = 1, nz
#endif
#ifdef __CUDA_DEVICE_FUNC__
call sincos_loop_1_kernelgen_blockidx_y(j, 1, ny)
#else
do j = 1, ny
#endif
#ifdef __CUDA_DEVICE_FUNC__
call sincos_loop_1_kernelgen_blockidx_x(i, 1, nx)
#else
do i = 1, nx
endif
  xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
endif
enddo
#endif
#ifdef __CUDA_DEVICE_FUNC__
enddo
#endif
#ifdef __CUDA_DEVICE_FUNC__
enddo
#endif
end subroutine sincos_loop_1_kernelgen
```

xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))

The body of original loop

# 3: LLVM IR for device code (1/2)

```
; ModuleID = 'sincos.sincos_loop_1_kernelgen.cuda.device.F90.ir'
target datalayout = "e-p:64:64:64-i1:8:8-i8:8:8-i16:16:16-i32:32:32-i64:64:64-f32:32:32-f64:64:64-
v64:64:64-v128:128:128-a0:0:64-s0:64:64-f80:128:128-f128:128:128-n8:16:32:64"
target triple = "x86_64-unknown-linux-gnu"

module asm "\09.ident\09\22GCC: (GNU) 4.5.4 20110810 (prerelease) LLVM: 136347M\22"

define void @sincos_loop_1_kernelgen_(i32* nocapture %nz, i32* nocapture %ny, i32* nocapture %nx, [0
x float]* %xy, [0 x float]* %x, [0 x float]* %y) nounwind uwtable {
entry:
    %memtmp = alloca i32, align 4
    %memtmp3 = alloca i32, align 4
    %memtmp4 = alloca i32, align 4
    %0 = load i32* %nx, align 4
    %1 = sext i32 %0 to i64
    %2 = icmp slt i64 %1, 0
    %3 = select i1 %2, i64 0, i64 %1
    %4 = load i32* %ny, align 4
    %5 = sext i32 %4 to i64
    %6 = mul nsw i64 %3, %5
    %7 = icmp slt i64 %6, 0
    %8 = select i1 %7, i64 0, i64 %6
    %not = xor i64 %3, -1
    %9 = sub nsw i64 %not, %8
    %10 = load i32* %nz, align 4
    call void (i32*, i32, i32, ...)* @sincos_loop_1_kernelgen_blockidx_z(i32* noalias %memtmp, i32 1,
i32 %10) nounwind
    %11 = load i32* %ny, align 4
    call void (i32*, i32, i32, ...)* @sincos_loop_1_kernelgen_blockidx_y(i32* noalias %memtmp3, i32 1,
i32 %11) nounwind
```

# 3: LLVM IR for device code (1/2)

```
; ModuleID = 'sincos.sincos_loop_1_kernelgen.cuda.device.F90.ir'
target datalayout = "e-p:64:64:64-i1:8:8-i8:8:8-i16:16:16-i32:32:32-i64:64:64-f32:32:32-f64:64:64-v64:64:64-v128:128:128-a0:0:64-s0:64:64-f80:128:128-f128:128:128-n8:16:32:64"
target triple = "x86_64-unknown-linux-gnu"
```

```
module asm "\09.ident\09\22GCC: (GNU) 4.5.4 20110810 (prerelease) LLVM: 136347M\22"
```

```
define void @sincos_loop_1_kernelgen_(i32* nocapture %nz, i32* nocapture %ny, i32* nocapture %nx, [0 x float]* %xy, [0 x float]* %x, [0 x float]* %y) nounwind uwtable {
```

entry:

subroutine sincos\_loop\_1\_kernelgen(nz, ny, nx, xy, x, y)

```
    %memtmp = alloca i32, align 4
    %memtmp4 = alloca i32, align 4
    %0 = load i32* %nx, align 4
    %1 = sext i32 %0 to i64
    %2 = icmp slt i64 %1, 0
    %3 = select i1 %2, i64 0, i64 %1
    %4 = load i32* %ny, align 4
    %5 = sext i32 %4 to i64
    %6 = mul nsw i64 %3, %5
    %7 = icmp slt i64 %6, 0
    %8 = select i1 %7, i64 0, i64 %6
    %not = xor i64 %3, -1
    %9 = sub nsw i64 %not, %8
    %10 = load i32* %nz, align 4
    call void (i32*, i32, i32, ...)* @sincos_loop_1_kernelgen_blockidx_z(i32* noalias %memtmp, i32 1, i32 %10) nounwind
    %11 = load i32* %ny, align 4
    call void (i32*, i32, i32, ...)* @sincos_loop_1_kernelgen_blockidx_y(i32* noalias %memtmp3, i32 1, i32 %11) nounwind
```

# 3: LLVM IR for device code (1/2)

```
; ModuleID = 'sincos.sincos_loop_1_kernelgen.cuda.device.F90.ir'
target datalayout = "e-p:64:64:64-i1:8:8-i8:8:8-i16:16:16-i32:32:32-i64:64:64-f32:32:32-f64:64:64-v64:64:64-v128:128:128-a0:0:64-s0:64:64-f80:128:128-f128:128:128-n8:16:32:64"
target triple = "x86_64-unknown-linux-gnu"

module asm "\09.ident\09\22GCC: (GNU) 4.5.4 20110810 (prerelease) LLVM: 136347M\22"

define void @sincos_loop_1_kernelgen_(i32* nocapture %nz, i32* nocapture %ny, i32* nocapture %nx, [0 x float]* %xy, [0 x float]* %x, [0 x float]* %y) nounwind uwtable {
entry:
  %memtmp = alloca i32, align 4
  %memtmp3 = alloca i32, align 4
  %memtmp4 = alloca i32, align 4
  %0 = load i32* %nx, align 4
  %1 = sext i32 %0 to i64
  %2 = icmp slt i64 %1, 0
  %3 = select i1 %2, i64 0, i64 %1
  %4 = load i32* %ny, align 4
  %5 = sext i32 %4 to i64
  %6 = mul nsw i64 %3, %5
  %7 = icmp slt i64 %6, 0
  %8 = select i1 %7, i64 0, i64 %6
  %not =
  %9 = su
  %10 =
    call sincos_loop_1_kernelgen_blockidx_z(k, 1, nz)
  call void (i32*, i32, i32, ...)* @sincos_loop_1_kernelgen_blockidx_z(i32* noalias %memtmp, i32 1, i32 %10) nounwind
  %11 = load i32* %ny, align 4
  call void (i32*, i32, i32, ...)* @sincos_loop_1_kernelgen_blockidx_y(i32* noalias %memtmp3, i32 1, i32 %11) nounwind
```

# 3: LLVM IR for device code (2/2)

```
%12 = load i32* %nx, align 4
call void (i32*, i32, i32, ...)* @sincos_loop_1_kernelgen_blockidx_x(i32* noalias %memtmp4, i32
1, i32 %12) nounwind
%13 = load i32* %memtmp4, align 4
%14 = sext i32 %13 to i64
%15 = load i32* %memtmp, align 4
%16 = sext i32 %15 to i64
%17 = mul nsw i64 %16, %8
%18 = load i32* %memtmp3, align 4
%19 = sext i32 %18 to i64
%20 = mul nsw i64 %19, %3
%21 = add i64 %14, %9
%22 = add i64 %21, %17
%23 = add i64 %22, %20
%24 = getelementptr [0 x float]* %x, i64 0, i64 %23
%25 = load float* %24, align 4
%26 = call float @sinf(float %25) nounwind readnone
%27 = getelementptr [0 x float]* %y, i64 0, i64 %23
%28 = load float* %27, align 4
%29 = call float @cosf(float %28) nounwind readnone
%30 = fadd float %26, %29
%31 = getelementptr [0 x float]* %xy, i64 0, i64 %23
store float %30, float* %31, align 4
ret void
}
```

```
declare void @sincos_loop_1_kernelgen_blockidx_z(i32* noalias, i32, i32, ...)
declare void @sincos_loop_1_kernelgen_blockidx_y(i32* noalias, i32, i32, ...)
declare void @sincos_loop_1_kernelgen_blockidx_x(i32* noalias, i32, i32, ...)
declare float @sinf(float) nounwind readnone
declare float @cosf(float) nounwind readnone
```

# 3: LLVM IR for device code (2/2)

```
%12 = load i32* %nx, align 4
call void (i32*, i32, i32, ...)* @sincos_loop_1_kernelgen_blockidx_x(i32* noalias %memtmp4, i32
1, i32 %12) nounwind
%13 = load i32* %memtmp4, align 4
%14 = sext i32 %13 to i64
%15 = load i32* %memtmp, align 4
%16 = sext i32 %15 to i64
%17 = mul nsw i64 %16, %8
%18 = load i32* %memtmp3, align 4
%19 = sext i32 %18 to i64
%20 = mul nsw i64 %19, %3
%21 = a
%22 = a
%23 = a
xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
%24 = getelementptr [0 x float]* %x, i64 0, i64 %23
%25 = load float* %24, align 4
%26 = call float @sinf(float %25) nounwind readnone
%27 = getelementptr [0 x float]* %y, i64 0, i64 %23
%28 = load float* %27, align 4
%29 = call float @cosf(float %28) nounwind readnone
%30 = fadd float %26, %29
%31 = getelementptr [0 x float]* %xy, i64 0, i64 %23
store float %30, float* %31, align 4
ret void
}
```

```
declare void @sincos_loop_1_kernelgen_blockidx_z(i32* noalias, i32, i32, ...)
declare void @sincos_loop_1_kernelgen_blockidx_y(i32* noalias, i32, i32, ...)
declare void @sincos_loop_1_kernelgen_blockidx_x(i32* noalias, i32, i32, ...)
declare float @sinf(float) nounwind readnone
declare float @cosf(float) nounwind readnone
```



# 4: C code for LLVM IR (1/3)

```
void sincos_loop_1_kernelgen_(
#ifdef __OPENCL_DEVICE_FUNC__
__global
#endif // __OPENCL_DEVICE_FUNC__
unsigned int *llvm_cbe_nz,
#ifdef __OPENCL_DEVICE_FUNC__
__global
#endif // __OPENCL_DEVICE_FUNC__
unsigned int *llvm_cbe_ny,
#ifdef __OPENCL_DEVICE_FUNC__
__global
#endif // __OPENCL_DEVICE_FUNC__
unsigned int *llvm_cbe_nx,
#ifdef __OPENCL_DEVICE_FUNC__
__global
#endif // __OPENCL_DEVICE_FUNC__
l_unnamed_0 (*llvm_cbe_xy),
#ifdef __OPENCL_DEVICE_FUNC__
__global
#endif // __OPENCL_DEVICE_FUNC__
l_unnamed_0 (*llvm_cbe_x),
#ifdef __OPENCL_DEVICE_FUNC__
__global
#endif // __OPENCL_DEVICE_FUNC__
l_unnamed_0 (*llvm_cbe_y)) {
    unsigned int llvm_cbe_memtmp;      /* Address-exposed local */
    unsigned int llvm_cbe_memtmp3;     /* Address-exposed local */
    unsigned int llvm_cbe_memtmp4;     /* Address-exposed local */
    unsigned int llvm_cbe_tmp_1;
    unsigned long long llvm_cbe_tmp_2;
    unsigned long long llvm_cbe_tmp_3;
```

# 4: C code for LLVM IR (1/3)

```
void sincos_loop_1_kernelgen (  
#ifdef __OPENCL_DEVICE_FUNC__  
__global  
#endif // __OPENCL_DEVICE_FUNC__  
unsigned int *llvm_cbe_nz,  
#ifdef __OPENCL_DEVICE_FUNC__  
__global  
#endif // __OPENCL_DEVICE_FUNC__  
unsigned int *llvm_cbe_ny,  
#ifdef __OPENCL_DEVICE_FUNC__  
__global  
#endif // __OPENCL_DEVICE_FUNC__  
unsigned int *llvm_cbe_nx,  
#ifdef __OPENCL_DEVICE_FUNC__  
__global  
#endif // __OPENCL_DEVICE_FUNC__  
l_unnamed_0 (*llvm_cbe_xy),  
#ifdef __OPENCL_DEVICE_FUNC__  
__global  
#endif // __OPENCL_DEVICE_FUNC__  
l_unnamed_0 (*llvm_cbe_x),  
#ifdef __OPENCL_DEVICE_FUNC__  
__global  
#endif // __OPENCL_DEVICE_FUNC__  
l_unnamed_0 (*llvm_cbe_y)) {  
    unsigned int llvm_cbe_memtmp; /* Address-exposed local */  
    unsigned int llvm_cbe_memtmp3; /* Address-exposed local */  
    unsigned int llvm_cbe_memtmp4; /* Address-exposed local */  
    unsigned int llvm_cbe_tmp_1;  
    unsigned long long llvm_cbe_tmp_2;  
    unsigned long long llvm_cbe_tmp_3;
```

In case of OpenCL, add \_\_global attribute to subroutine arguments

```
    unsigned int llvm_cbe_memtmp; /* Address-exposed local */  
    unsigned int llvm_cbe_memtmp3; /* Address-exposed local */  
    unsigned int llvm_cbe_memtmp4; /* Address-exposed local */  
    unsigned int llvm_cbe_tmp_1;  
    unsigned long long llvm_cbe_tmp_2;  
    unsigned long long llvm_cbe_tmp_3;
```

# 4: C code for LLVM IR (2/3)

```
unsigned int llvm_cbe_tmp__4;  
unsigned long long llvm_cbe_tmp__5;  
unsigned long long llvm_cbe_tmp__6;  
unsigned int llvm_cbe_tmp__7;  
unsigned int llvm_cbe_tmp__8;  
unsigned int llvm_cbe_tmp__9;  
unsigned int llvm_cbe_tmp__10;  
unsigned int llvm_cbe_tmp__11;  
unsigned int llvm_cbe_tmp__12;  
unsigned long long llvm_cbe_tmp__13;  
float llvm_cbe_tmp__14;  
float llvm_cbe_tmp__15;  
float llvm_cbe_tmp__16;  
float llvm_cbe_tmp__17;  
  
llvm_cbe_tmp__1 = *llvm_cbe_nx;  
llvm_cbe_tmp__2 = ((signed long long )(signed int )llvm_cbe_tmp__1);  
llvm_cbe_tmp__3 = (((((signed long long )llvm_cbe_tmp__2) < ((signed long long )0ull))) ?  
(0ull) : (llvm_cbe_tmp__2));  
llvm_cbe_tmp__4 = *llvm_cbe_ny;  
llvm_cbe_tmp__5 = (((unsigned long long )(((unsigned long long )llvm_cbe_tmp__3) * ((unsigned long  
long )(((signed long long )(signed int )llvm_cbe_tmp__4)))));  
llvm_cbe_tmp__6 = (((((signed long long )llvm_cbe_tmp__5) < ((signed long long )0ull))) ?  
(0ull) : (llvm_cbe_tmp__5));  
llvm_cbe_tmp__7 = *llvm_cbe_nz;  
sincos_loop_1_kernelgen_blockidx_z((&llvm_cbe_memtmp), 1u, llvm_cbe_tmp__7);  
llvm_cbe_tmp__8 = *llvm_cbe_ny;  
sincos_loop_1_kernelgen_blockidx_y((&llvm_cbe_memtmp3), 1u, llvm_cbe_tmp__8);  
llvm_cbe_tmp__9 = *llvm_cbe_nx;  
sincos_loop_1_kernelgen_blockidx_x((&llvm_cbe_memtmp4), 1u, llvm_cbe_tmp__9);
```

# 4: C code for LLVM IR (2/3)

```
unsigned int llvm_cbe_tmp_4;  
unsigned long long llvm_cbe_tmp_5;  
unsigned long long llvm_cbe_tmp_6;  
unsigned int llvm_cbe_tmp_7;  
unsigned int llvm_cbe_tmp_8;  
unsigned int llvm_cbe_tmp_9;  
unsigned int llvm_cbe_tmp_10;  
unsigned int llvm_cbe_tmp_11;  
unsigned int llvm_cbe_tmp_12;  
unsigned long long llvm_cbe_tmp_13;  
float llvm_cbe_tmp_14;  
float llvm_cbe_tmp_15;  
float llvm_cbe_tmp_16;  
float llvm_cbe_tmp_17;  
  
llvm_cbe_tmp_1 = *llvm_cbe_nx;  
llvm_cbe_tmp_2 = ((signed long long )(signed int )llvm_cbe_tmp_1);  
llvm_cbe_tmp_3 = (((((signed long long )llvm_cbe_tmp_2) < ((signed long long )0ull))) ?  
(0ull) : (llvm_cbe_tmp_2));  
llvm_cbe_tmp_4 = *llvm_cbe_ny;  
llvm_cbe_tmp_5 = (((unsigned long long )(((unsigned long long )llvm_cbe_tmp_3) * ((unsigned long  
long )(((signed long long )(signed int )llvm_cbe_tmp_4))))));  
llvm_cbe_tmp_6 = (((((signed long long )llvm_cbe_tmp_5) < ((signed long long )0ull))) ?  
(0ull) : (llvm_cbe_tmp_5));  
call sincos_loop_1_kernelgen_blockidx_z(k, 1, nz)  
sincos_loop_1_kernelgen_blockidx_z((&llvm_cbe_memtmp), 1u, llvm_cbe_tmp_7);  
llvm_cbe_tmp_8 = *llvm_cbe_ny;  
sincos_loop_1_kernelgen_blockidx_y((&llvm_cbe_memtmp3), 1u, llvm_cbe_tmp_8);  
llvm_cbe_tmp_9 = *llvm_cbe_nx;  
sincos_loop_1_kernelgen_blockidx_x((&llvm_cbe_memtmp4), 1u, llvm_cbe_tmp_9);
```

# 4: C code for LLVM IR (3/3)

```
llvm_cbe_tmp__11 = *(&llvm_cbe_memtmp);
llvm_cbe_tmp__12 = *(&llvm_cbe_memtmp3);
llvm_cbe_tmp__13 = (((unsigned long long )(((unsigned long long )(((unsigned long long )
(((unsigned long long )(((unsigned long long )(((unsigned long long )(((signed long long )(signed
int )llvm_cbe_tmp__10)))) + ((unsigned long long )(((unsigned long long )(((unsigned long long )
(llvm_cbe_tmp__3 ^ 18446744073709551615ull)) - ((unsigned long long )llvm_cbe_tmp__6))))))))) +
((unsigned long long )(((unsigned long long )(((unsigned long long )(((signed long long )(signed
int )llvm_cbe_tmp__11))) * ((unsigned long long )llvm_cbe_tmp__6))))))))) + ((unsigned long long )
(((unsigned long long )(((unsigned long long )(((unsigned long long )(((signed long long )(signed
int )llvm_cbe_tmp__12))) * ((unsigned long long )llvm_cbe_tmp__6))))))))) + ((unsigned long long )
* ((unsigned long long )llvm_cbe_tmp__12)))
xy(i, j, k) = sin(x(i, j, k)) + cos(y(i, j, k))
llvm_cbe_tmp__15 = sinf(llvm_cbe_tmp__14);
llvm_cbe_tmp__16 = *(&(*llvm_cbe_y).array[(((signed long long )llvm_cbe_tmp__13))]);
llvm_cbe_tmp__17 = cosf(llvm_cbe_tmp__16);
*(&(*llvm_cbe_xy).array[(((signed long long )llvm_cbe_tmp__13))]) = (((float )(llvm_cbe_tmp__15 +
llvm_cbe_tmp__17)));
return;
}
```

# Compiling sincos

```
[marcusmae@noisy ~]$ cd Programming/kernelgen/trunk/tests/performance/sincos/
[marcusmae@noisy sincos]$ make 32/sincos
kgen-gfortran -Wk,--host-compiler=/usr/bin/gfortran
-I/home/marcusmae/Programming/kernelgen/trunk/include -Wk,--cpu-compiler=gcc
-Wk,--cuda-compiler=nvcc -m32 -Wk,--opencl-compiler=kgen-opencl-embed -Wk,--
kernel-target=cpu,cuda,opencl, -O3 -g -c sincos.f90 -o 32/sincos.o
kernelgen    >> sincos.f90:42: portable 3-dimensional loop
kernelgen    >> sincos.f90:42: selecting this loop
c    >> ptxas info      : Compiling entry function 'sincos_loop_1_kernelgen_cuda'
for 'sm_20'
c    >> ptxas info      : Function properties for sincos_loop_1_kernelgen_cuda
c    >>      56 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
c    >> ptxas info      : Used 12 registers, 4+0 bytes lmem, 56 bytes cmem[0], 24
bytes cmem[2], 44 bytes cmem[16]
/usr/bin/gcc -I/home/marcusmae/Programming/kernelgen/trunk/include
-I/home/marcusmae/opt/kgen/include -m32 -O3 -g -std=c99 -I/opt/kgen/include -c
main.c -o 32/main.o
kgen-gfortran -Wk,--host-compiler=/usr/bin/gfortran
-I/home/marcusmae/Programming/kernelgen/trunk/include -Wk,--cpu-compiler=gcc
-Wk,--cuda-compiler=nvcc -m32 -Wk,--opencl-compiler=kgen-opencl-embed -Wk,--
kernel-target=cpu,cuda,opencl, 32/main.o 32/sincos.o -o 32/sincos
```

# Compiling sincos

```
[marcusmae@noisy ~]$ cd Programming/kernelgen/trunk/tests/performance/sincos/  
[marcusmae@noisy sincos]$ make 32/sincos
```

```
kgen-gfortran -Wk,--host-compiler=/usr/bin/gfortran  
-I/home/marcusmae/Programming/kernelgen/trunk/include -Wk,--cpu-compiler=gcc  
-Wk,--cuda-compiler=nvcc -m32 -Wk,--opencl-compiler=kgen-opencl-embed -Wk,--  
kernel-target=cpu,cuda,opencl, -O3 -g -c sincos.f90 -o 32/sincos.o
```

kernelgen  
kernelgen  
KernelGen compilation command, specifying target  
devices and compilers to use

```
c  >> ptxas info      : Function properties for sincos_loop_1_kernelgen_cuda  
for 'sm_20'  
c  >> ptxas info      : 56 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads  
c  >> ptxas info      : Used 12 registers, 4+0 bytes lmem, 56 bytes cmem[0], 24  
bytes cmem[2], 44 bytes cmem[16]  
/usr/bin/gcc -I/home/marcusmae/Programming/kernelgen/trunk/include  
-I/home/marcusmae/opt/kgen/include -m32 -O3 -g -std=c99 -I/opt/kgen/include -c  
main.c -o 32/main.o  
kgen-gfortran -Wk,--host-compiler=/usr/bin/gfortran  
-I/home/marcusmae/Programming/kernelgen/trunk/include -Wk,--cpu-compiler=gcc  
-Wk,--cuda-compiler=nvcc -m32 -Wk,--opencl-compiler=kgen-opencl-embed -Wk,--  
kernel-target=cpu,cuda,opencl, 32/main.o 32/sincos.o -o 32/sincos
```

# Compiling sincos

```
[marcusmae@noisy ~]$ cd Programming/kernelgen/trunk/tests/performance/sincos/
[marcusmae@noisy sincos]$ make 32/sincos
kgen-gfortran -Wk,--host-compiler=/usr/bin/gfortran
-I/home/marcusmae/Programming/kernelgen/trunk/include -Wk,--cpu-compiler=gcc
-Wk,--cuda-compiler=nvcc -m32 -Wk,--opencl-compiler=kgen-opencl-embed -Wk,--
kernel-target=cpu,cuda,opencl, -03 -g -c sincos.f90 -o 32/sincos.o
```

```
kernelgen  >> sincos.f90:42: portable 3-dimensional loop
kernelgen  >> sincos.f90:42: selecting this loop
```

```
c  >> p KernelGen reports indentified portable loops and those of them selected to have device version
for 'sm_
c  >> p
c  >>      56 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
c  >> ptxas info      : Used 12 registers, 4+0 bytes lmem, 56 bytes cmem[0], 24
bytes cmem[2], 44 bytes cmem[16]
/usr/bin/gcc -I/home/marcusmae/Programming/kernelgen/trunk/include
-I/home/marcusmae/opt/kgen/include -m32 -03 -g -std=c99 -I/opt/kgen/include -c
main.c -o 32/main.o
kgen-gfortran -Wk,--host-compiler=/usr/bin/gfortran
-I/home/marcusmae/Programming/kernelgen/trunk/include -Wk,--cpu-compiler=gcc
-Wk,--cuda-compiler=nvcc -m32 -Wk,--opencl-compiler=kgen-opencl-embed -Wk,--
kernel-target=cpu,cuda,opencl, 32/main.o 32/sincos.o -o 32/sincos
```



# Compiling sincos

```
[marcusmae@noisy ~]$ cd Programming/kernelgen/trunk/tests/performance/sincos/
[marcusmae@noisy sincos]$ make 32/sincos
kgen-gfortran -Wk,--host-compiler=/usr/bin/gfortran
-I/home/marcusmae/Programming/kernelgen/trunk/include -Wk,--cpu-compiler=gcc
-Wk,--cuda-compiler=nvcc -m32 -Wk,--opencl-compiler=kgen-opencl-embed -Wk,--
kernel-target=cpu,cuda,opencl, -O3 -g -c sincos.f90 -o 32/sincos.o
kernelgen    >> sincos.f90:42: portable 3-dimensional loop
kernelgen    >> sincos.f90:42: selecting this loop
```

```
c    >> ptxas info      : Compiling entry function 'sincos_loop_1_kernelgen_cuda'
for 'sm_20'
c    >> ptxas info      : Function properties for sincos_loop_1_kernelgen_cuda
c    >>      56 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
c    >> ptxas info      : Used 12 registers, 4+0 bytes lmem, 56 bytes cmem[0], 24
bytes cmem[2], 44 bytes cmem[16]
```

```
/usr/bin/gcc -I/home/marcusmae/Programming/kernelgen/trunk/include
-I/home/marcusmae/opt/kgen/trunk/include -I/opt/kgen/include -c
main.c -o 32/main.o
kgen-gfortran -Wk,--host-compiler=/usr/bin/gfortran
-I/home/marcusmae/Programming/kernelgen/trunk/include -Wk,--cpu-compiler=gcc
-Wk,--cuda-compiler=nvcc -m32 -Wk,--opencl-compiler=kgen-opencl-embed -Wk,--
kernel-target=cpu,cuda,opencl, 32/main.o 32/sincos.o -o 32/sincos
```

Output from ptx-as

# Compiling sincos

```
[marcusmae@noisy ~]$ cd Programming/kernelgen/trunk/tests/performance/sincos/
[marcusmae@noisy sincos]$ make 32/sincos
kgen-gfortran -Wk,--host-compiler=/usr/bin/gfortran
-I/home/marcusmae/Programming/kernelgen/trunk/include -Wk,--cpu-compiler=gcc
-Wk,--cuda-compiler=nvcc -m32 -Wk,--opencl-compiler=kgen-opencl-embed -Wk,--
kernel-target=cpu,cuda,opencl, -O3 -g -c sincos.f90 -o 32/sincos.o
kernelgen    >> sincos.f90:42: portable 3-dimensional loop
kernelgen    >> sincos.f90:42: selecting this loop
c    >> ptxas info      : Compiling entry function 'sincos_loop_1_kernelgen_cuda'
for 'sm_20'
c    >> ptxas info      : Function properties for sincos_loop_1_kernelgen_cuda
c    >>      56 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
c    >> ptxas info      : Used 12 registers, 4+0 bytes lmem, 56 bytes cmem[0], 24
bytes cmem[2], 44 bytes cmem[16]
/usr/bin/gcc -I/home/marcusmae/Programming/kernelgen/trunk/include
-I/home/marcusmae/opt/kgen/ -std=c99 -I/opt/kgen/include -c
main.c -o 32/main.o
kgen-gfortran -Wk,--host-compiler=/usr/bin/gfortran
-I/home/marcusmae/Programming/kernelgen/trunk/include -Wk,--cpu-compiler=gcc
-Wk,--cuda-compiler=nvcc -m32 -Wk,--opencl-compiler=kgen-opencl-embed -Wk,--
kernel-target=cpu,cuda,opencl, 32/main.o 32/sincos.o -o 32/sincos
```

Linker command

# Testing sincos

## **# By default – execute on CPU**

```
[marcusmae@noisy sincos]$ 64/sincos 128 128 4 4 4 16  
gpu time = 0.663560 sec  
cpu time = 0.642710 sec  
max diff = 0.000000e+00 @ 0
```

## **# Set default runmode to 2 to execute CUDA versions of all kernels**

```
[marcusmae@noisy sincos]$ kernelgen_runmode=2 sincos_loop_1_kernelgen=2  
64/sincos 128 128 4 4 4 16  
gpu time = 0.369378 sec  
cpu time = 0.643688 sec  
max diff = 2.384186e-07 @ 310
```

## **# Set default runmode to 4 to execute OpenCL versions of all kernels**

```
[marcusmae@noisy sincos]$ kernelgen_runmode=4 sincos_loop_1_kernelgen=4  
32/sincos 128 128 4 4 4 16  
gpu time = 0.443252 sec  
cpu time = 1.134842 sec  
max diff = 2.384186e-07 @ 4349389
```

# Testing sincos


## # Add debug output filter bits to show more info

```
[marcusmae@noisy sincos]$ kernelgen_debug_output=11 kernelgen_runmode=2 kgen/32/sincos 512 512 64 1 1 1
launch.c:70 kernelgen message (debug) Launching sincos_loop_1_kernelgen_cuda for device NVIDIA Corporation:0
runmode "cuda"
parse_args.h:69 kernelgen message (debug) arg "unknown" ref = 0xffff9837c, size = 4, desc = 0xffff9837c
parse_args.h:69 kernelgen message (debug) arg "unknown" ref = 0xffff98378, size = 4, desc = 0xffff98378
parse_args.h:69 kernelgen message (debug) arg "unknown" ref = 0xffff98374, size = 4, desc = 0xffff98374
parse_args.h:69 kernelgen message (debug) arg "unknown" ref = 0xe3443008, size = 67108864, desc = 0xe3443008
parse_args.h:69 kernelgen message (debug) arg "unknown" ref = 0xf3447008, size = 67108864, desc = 0xf3447008
parse_args.h:69 kernelgen message (debug) arg "unknown" ref = 0xeb445008, size = 67108864, desc = 0xeb445008
map_cuda.h:107 kernelgen message (debug) symbol "unknown" maps memory segment [0xffff9837c .. 0xffff98380] to
[0x5400000 .. 0x5400004]
map_cuda.h:107 kernelgen message (debug) symbol "unknown" maps memory segment [0xffff98378 .. 0xffff9837c] to
[0x5400200 .. 0x5400204]
map_cuda.h:107 kernelgen message (debug) symbol "unknown" maps memory segment [0xffff98374 .. 0xffff98378] to
[0x5400400 .. 0x5400404]
map_cuda.h:107 kernelgen message (debug) symbol "unknown" maps memory segment [0xe3443008 .. 0xe7443008] to
[0x5500000 .. 0x9500000]
map_cuda.h:107 kernelgen message (debug) symbol "unknown" maps memory segment [0xf3447008 .. 0xf7447008] to
[0x9500000 .. 0xd500000]
map_cuda.h:107 kernelgen message (debug) symbol "unknown" maps memory segment [0xeb445008 .. 0xef445008] to
[0xd500000 .. 0x11500000]
gpu time = 0.373937 sec
cpu time = 1.136829 sec
max diff = 1.192093e-07 @ 17
```

## **4. Testing unoptimized generator**


# Example: sincos - performance

Performance of CPU binary generated by gfortran and CUDA kernel by KernelGen compared to host and device perfs, using PGI Accelerator 11.8 (orange)




# Example: sincos - performance

Performance of CPU binary generated by gfortran and CUDA kernel by KernelGen compared to host and device perfs, using PGI Accelerator 11.8 (orange)



# Example: sincos - performance


Performance of CPU binary generated by gfortran and CUDA kernel by KernelGen compared to host and device perfs, using PGI Accelerator 11.8 (orange)






# Example: sincos - performance

Performance of CPU binary generated by gfortran and CUDA kernel by KernelGen compared to host and device perfs, using PGI Accelerator 11.8 (orange)




# Example: sincos - performance


Performance of CPU binary generated by gfortran and OpenCL/CUDA kernels by KernelGen, compared to host and device perfs using PGI Accelerator 11.8 (orange)




# COSMO - coverage




# COSMO - coverage




# COSMO - coverage




# COSMO - coverage



# COSMO - coverage




# COSMO - coverage






# COSMO - coverage




# COSMO - performance

The generated kernel speedup, times




The generated kernel name

# WRF - coverage



# WRF - performance



## The generated kernel name

# Why slowdown?

- KernelGen does not **yet** utilize multiple threads inside thread blocks

**# Performance of CUDA version with 4x4x16 threads grid and copy-in/copy-out data masking is very similar to PGI's**


```
[marcusmae@noisy sincos]$ 64/sincos 128 128 4 4 4 16 0  
gpu time = 0.163355 sec  
cpu time = 0.641843 sec  
max diff = 2.384186e-07 @ 310
```

**# Performance of CUDA version with 1x1x1 threads grid and full data copying is very similar to KernelGen's**

```
[marcusmae@noisy sincos]$ cuda/64/sincos 512 512 64 1 1 1 1  
gpu time = 0.450221 sec  
cpu time = 0.642180 sec  
max diff = 2.384186e-07 @ 310
```

# Why slowdown?

- KernelGen does not **yet** utilize multiple threads inside thread blocks
- Threads in blocks would be possible with **tiling** optimization implemented (from LLVM/Polly)



## **5. Development schedule**

# Stage 1 (April - June)



- Put together all necessary toolchain parts, write the main script
- Test C code generation, file bugs to llvm, patch C backend for CUDA support
- Complete existing host-device code split transform (previously started in 2009 for CellBE)
- Implement kernel invocation runtime
- Implement kernel self-checking runtime
- Compile COSMO with toolchain and present charts showing the percentage of successfully generated kernels with checked correct results



# Stage 2 (July - October)

- Improve support/coverage
  - More testing on COSMO and other models, file bugs (+2 RHM fellows)
  - Fix the most hot bugs in host-device code split transform
  - Use Polly/Pluto for threading and more accurate capable loops recognition
  - Support link-time generation for kernels with external dependencies
- Improve efficiency
  - Use shared memory in stencils (+1 contractor)
  - Implement both zero-copy and active data synchronization modes
  - Kernel invocation configs caching
  - [variant] Consider putting serial code into single GPU thread as well, to have the whole model instance running on GPU
  - [variant] Consider selective/prioritized data synchronization support, using data dependencies lookup
  - [variant, suggested by S.K.] CPU ↔ GPU work sharing inside MPI process
- Compare performance with other generation tools
- Present the work and carefully listen to feedback

# Stage 2 (July - October)

- Improve support/coverage  – done  – in progress now
  - **More testing on COSMO and other models, file bugs (+2 RHM fellows)**
  - **Fix the most hot bugs in host-device code split transform**
  - **Use Polly/Pluto for threading and more accurate capable loops recognition**
  - **Support link-time generation for kernels with external dependencies**
- Improve efficiency
  - **Use shared memory in stencils (+1 contractor)**
  - **Implement both zero-copy and active data synchronization modes**
  - **Kernel invocation configs caching**
  - [variant] Consider putting serial code into single GPU thread as well, to have the whole model instance running on GPU
  - [variant] Consider selective/prioritized data synchronization support, using data dependencies lookup
  - [variant, suggested by S.K.] CPU ↔ GPU work sharing inside MPI process
- Compare performance with other generation tools
- Present the work and carefully listen to feedback

## **6. Team & resources**

# Team



Artem Petrov

(testing, coordination)

Dr Yulia Martynova

(WRF testing)

# Team



Artem Petrov

(testing, coordination)

Dr Yulia Martynova

(WRF testing)

---

Alexander Myltsev

(development, testing)

Dmitry Mikushin

(development, planning)

# Team



Artem Petrov

(testing, coordination)

Dr Yulia Martynova

(WRF testing)

---

Alexander Myltsev

(development, testing)

Dmitry Mikushin

(development, planning)

---

Support from  
communities:



LLVM



Polly/LLVM



gcc/gfortran

# Other projects used

- **g95-xml** – the XML markup for Fortran 95 source code based on g95 compiler (by Philippe Marguinaud). Used as input for code split transformations
- **LLVM Dragonegg** – bridge to utilize GCC as frontend to LLVM  $\Rightarrow$  compile Fortran code (by Duncan Sands et al)
- **LLVM C backend** – C code generator out of LLVM IR (by Chris Lattner, Duncan Sands et al)

# KernelGen preview release

Project source code, docs and binaries at HPCForge:

<http://hpcforge.org/projects/kernelgen/>

Binaries for 64-bit Fedora 15:

[kernelgen-0.1-cuda.x86\\_64.rpm](#)

[kernelgen-0.1-opencl.x86\\_64.rpm](#)

Documentation on wiki:

[Running the public test suite](#)

[Compiling \(for developers\)](#)



# Collaboration

## We provide:

- Source code and binaries
- User support, updates and bug fixes

## We need:

- Users feedback, testing and filing bugs
- Access to actual benchmarks (our COSMO is v4.13)
- Developers are welcome, especially skilled in LLVM and/or models

**Thank you! 😊 Questions?**