

Current Situation of the COSMO-Model

OR: Why has the COSMO-Model problems with new architectures?

- ➔ It only uses a pure MPI parallelization. Within one MPI task it is written for a flat memory access.
 - ➔ NEC SX-9: 13 % of peak
 - ➔ IBM pwr6: about 5-6 % of peak
 - ➔ Cray XT4: about 3-4 % of peak (?)
- ➔ One part is becoming more and more serious: I/O
 - ➔ many global communications involved
 - ➔ disk access is purely sequential
- ➔ Although the communications besides I/O are almost all local, the speedup degrades when using many processors.